

CS 33:

Computer Organization

Glenn Reinman
4731G Boelter Hall
reinman@cs.ucla.edu

TAs:

Sharath Gopal

Garett Ridge

Olivera Grujic

Marco Vitanza

Long Nguyen

Jason Zheng

Course Components

- Lectures
 - Higher level concepts
- Discussions
 - Applied concepts, important tools and skills for labs, clarification of lectures, exam coverage
- Labs
 - The heart of the course
 - Provide in-depth understanding of an aspect of systems
 - Programming and measurement

More Info

- Web

- Class web page hosted by CourseWeb
- Copies of lectures, assignments, exams, solutions
- Forum

- Office Hours

- Textbook

- Randal E. Bryant and David R. O'Hallaron. “Computer Systems: A Programmer’s Perspective”, **2nd Edition**, Prentice Hall 2010.

Grading

- Exams (45%)
 - Two in class exams (15% each)
 - Final (15%)
 - All exams are open book/open notes.
- Labs (50%)
 - 5 labs (10% each)
 - You must work alone on all labs
- Homework (5%)
 - 5 assignments (1% each)
 - Electronic submission only

Tentative Calendar

Week	M	T	W	R	F
1		Intro + Labs (1)		Bits and Bytes (2)	
2		Integers (2)		Floating Point (2)	Data Lab Due
3	HW #1 Due	Machine-Level Rep (3)		Machine-Level Rep (3)	
4	HW #2 Due	Machine-Level Rep (3)		Exam #1	Bomb Lab Due
5	HW #3 Due	Code Optimization (5)		Code Optimization (5)	
6		Memory (6)		Concurrency (12+handouts)	Buffer Lab Due
7	HW #4 Due	Concurrency (12+handouts)		Exam #2	
8		Concurrency (12+handouts)		I/O (10)	OpenMP Lab Due
9	HW #5 Due	Virtual Memory (9)		Linking + Exceptions (7,8)	
10		MIPS (handouts)		Review	CUDA Lab Due

- Homework and Labs Due via CourseWeb by Midnight

Cheating

- What is cheating?
 - Sharing code: either by copying, retyping, looking at, or supplying a copy of a file.
- What is NOT cheating?
 - Helping others use systems or tools.
 - Helping others with high-level design issues.
 - Helping others debug their code.
- Penalty for cheating:
 - At the discretion of the Associate Dean

Lab Facilities

- SEAS Administered Linux Machine

- lnxsrv02.seas.ucl.ac.uk
- Remote access only
 - Use ssh to log in with your SEAS account
- Please direct any account issues to the SEAS help desk as they are the only ones with root access on this machine

- Alternatives (Not Recommended)

- You may use other alternatives to develop your code
- **BUT: We will test on the SEAS machines**
 - **Your code must work correctly on these machines for credit**

Course Theme

- Abstraction is good, but don't forget reality!
- Abstractions have limits
 - Things are more complex in hardware than they look in C/Java!!
 - Bugs are hard to track/understand if looking only from a high-level point of view
- Useful outcomes
 - Become more effective programmers
 - Able to find and eliminate bugs efficiently
 - Able to tune program performance
 - Prepare for later “systems” classes in CS
 - Compilers, Operating Systems, Networks, Computer Architecture, Parallel Programming

The Compilation System

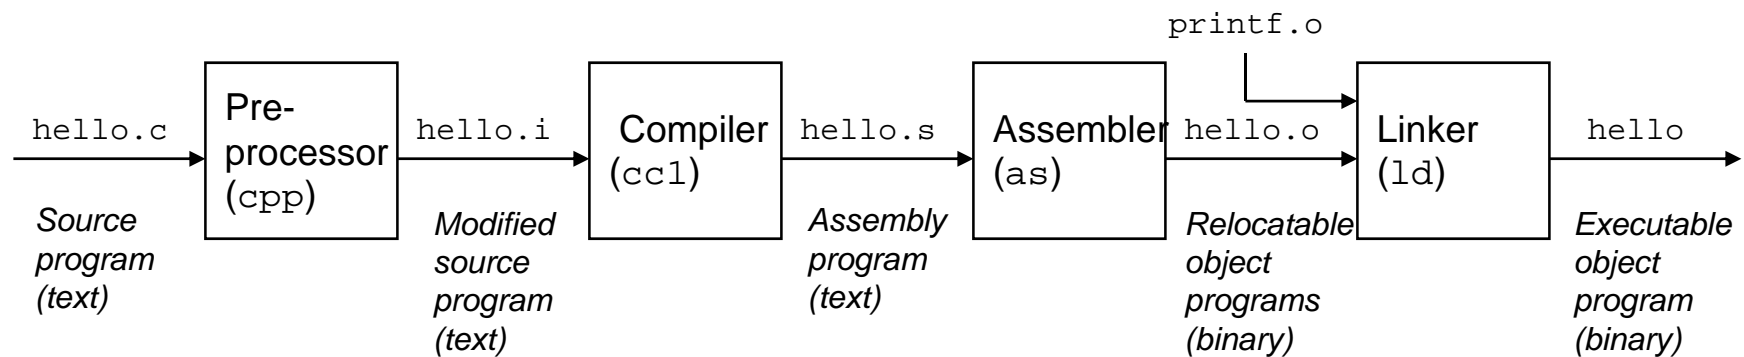
```
#include <stdio.h>
```

```
int main()
```

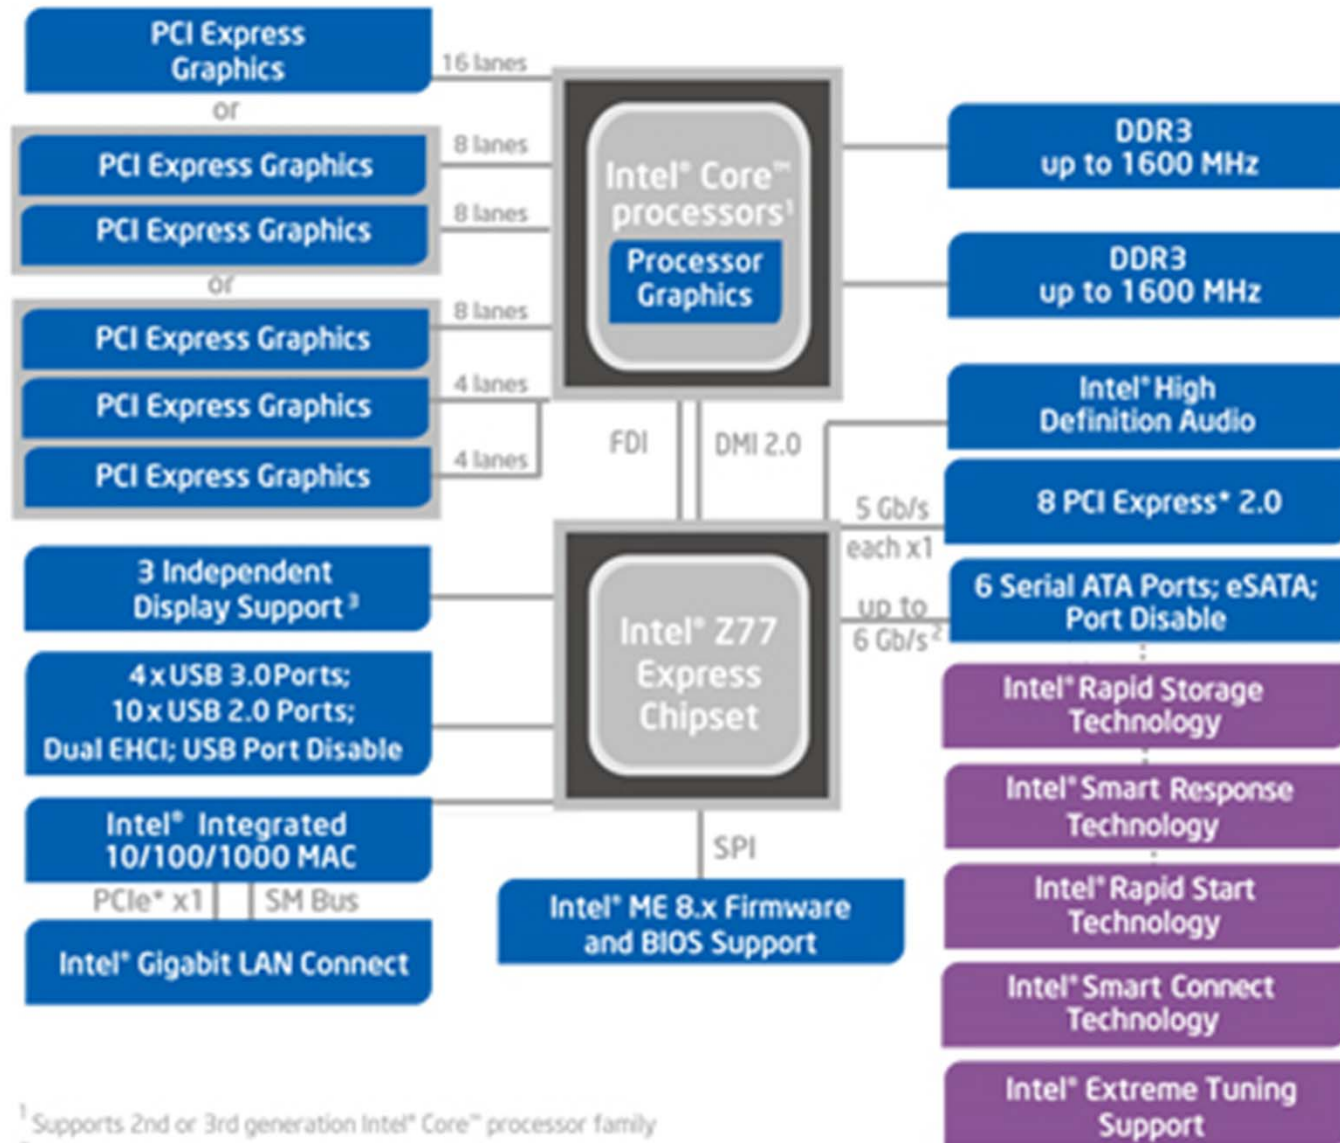
```
{
```

```
    printf("hello, world\n");
```

```
}
```



Sandy Bridge Architecture



¹ Supports 2nd or 3rd generation Intel® Core™ processor family

² All SATA ports capable of 3 Gb/s. 2 ports capable of 6 Gb/s

³ Available with 3rd generation Intel® Core™ processor

Great Reality #1

- *Computers do more than execute programs*
- They need to get data in and out
 - I/O system critical to program reliability and performance
- They communicate with each other over networks
 - Many system-level issues arise in presence of network
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross platform compatibility
 - Complex performance issues

Great Reality #2

- *Parallelism is key to future performance*
- Power-Efficient Performance comes from CMPs
 - Chip Multiprocessors (CMPs)
 - Multiple processor cores integrated onto a single silicon die or multichip module
 - Examples?
 - Intel's Core iX
 - Sun's Niagara 2
 - IBM/Sony's Cell
 - nVidia's Tesla, Fermi, Kepler
- Programmers need to learn how to exploit parallelism in their applications

Great Reality #3

- *You've got to know assembly & the machine!*
- Although you may never **program** in assembly
 - Compilers are much better & more patient than you are
- Assembly key to machine-level execution model
 - Behavior of programs in presence of bugs
 - High-level language model breaks down
 - Tuning program performance
 - Understanding sources of program inefficiency
 - Implementing system software
 - Compiler has machine code as target
 - Operating systems must manage process state

Great Reality #4

- *Memory Matters*
- Memory is not unbounded
 - It must be allocated and managed
 - Many applications are memory dominated
- Memory performance is not uniform
 - Cache and virtual memory effects can greatly affect program performance
 - Adapting program to characteristics of memory system can lead to major speed improvements
- Memory referencing bugs especially pernicious
 - Effects are distant in both time and space

Great Reality #5

- *There's more to performance than asymptotic complexity*
- Constant factors matter too!
 - 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- Must understand system to optimize performance
 - How programs are compiled and executed
 - Measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Great Reality #6

- *Int's are not Integers, Float's are not Reals*

- Examples

- Is $x^2 \geq 0$?

- IEEE float: Yes!

- 32 bit int:

- $40000 * 40000 \rightarrow 1600000000$

- $50000 * 50000 \rightarrow ??$

- Is $(x + y) + z = x + (y + z)$?

- Unsigned & Signed Int's: Yes!

- Float's:

- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$

- $1e20 + (-1e20 + 3.14) \rightarrow ??$

Computer Arithmetic

- Does not generate random values
 - Arithmetic operations have important mathematical properties
- Cannot assume “usual” properties
 - Due to finiteness of representations
 - Integer operations satisfy “ring” properties
 - Commutativity, associativity, distributivity
 - Floating point operations satisfy “ordering” properties
 - Monotonicity, values of signs
- Observation
 - Need to understand which abstractions apply in which contexts
 - Important issues for compiler writers and numerical application programmers

Tools and Skills

- C Programming
- x86 Assembly
- Basic CUDA
- Basic MIPS
- Debugging
 - gdb – GNU Debugger
- Rudimentary Understanding of
 - An Editor (i.e. emacs, vi)
 - Linux