


Machine-Level Programming: Procedures

Chapter 3 of B&O

Some notes adopted from Bryant and O'Hallaron

IA32 Stack


- Region of memory managed with stack discipline
- Grows toward lower addresses
- Register `%esp` indicates lowest stack address
 - address of top element



IA32 Stack Pushing

- Pushing


- `pushl Src`
- Fetch operand at *Src*
- Decrement `%esp` by 4
- Write operand at address given by `%esp`



IA32 Stack Popping

- **Popping**

- `popl Dest`
- Read operand at address given by `%esp`
- Increment `%esp` by 4
- Write to *Dest*



Procedure Control Flow


- Use stack to support procedure call and return
- Procedure call:
 - call *label*: Push return address on stack; Jump to *label*
- Return address value
 - Address of instruction beyond call
 - Example from disassembly
 - 804854e: e8 3d 06 00 00 call 8048b90 <main>
 - 8048553: 50 pushl %eax
 - Return address = 0x8048553
- Procedure return:
 - ret Pop address from stack; Jump to address

Procedure Call Example

804854e: e8 3d 06 00 00
8048553: 50


call 8048b90 <main>
pushl %eax

call 8048b90



%esp 0x108

%eip 0x804854e



%esp 0x104

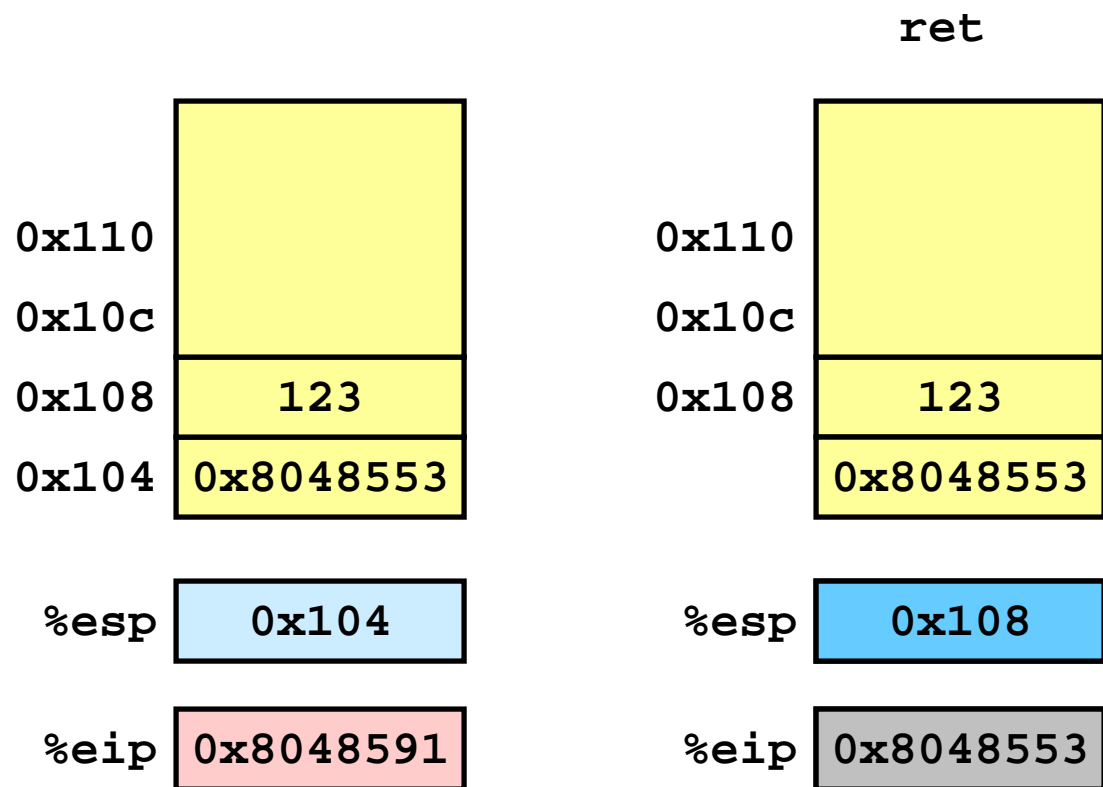
%eip 0x8048b90

%eip is program counter

Procedure Return Example

8048591: c3

ret



`%eip` is program counter

Stack-Based Languages

- Languages that Support Recursion
 - e.g., C, Pascal, Java
 - Code must be “Reentrant”
 - Multiple simultaneous instantiations of single procedure
 - Need some place to store state of each instantiation
 - Arguments
 - Local variables
 - Return pointer
- Stack Discipline
 - State for given procedure needed for limited time
 - From when called to when return
 - Callee returns before caller does
- Stack Allocated in Frames
 - state for single procedure instantiation

Call Chain Example

Code Structure


```
yoo(...)  
{  
  •  
  •  
  who( );  
  •  
  •  
}
```

```
who(...)  
{  
  • • •  
  amI( );  
  • • •  
  amI( );  
  • • •  
}
```

```
amI(...)  
{  
  •  
  •  
  amI( );  
  •  
  •  
}
```

- Procedure amI recursive

Call Chain



Stack Frames

- **Contents**

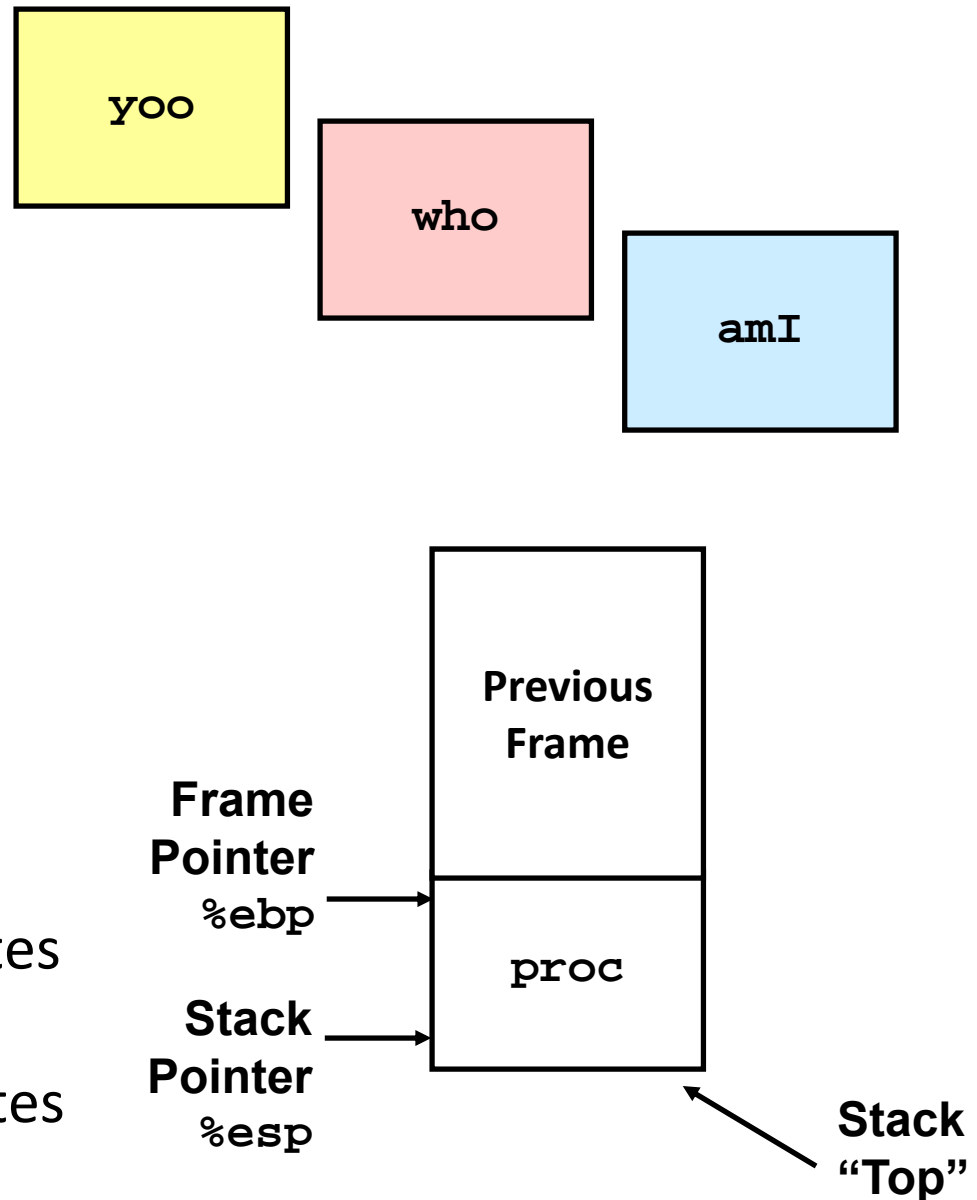
- Local variables
- Return information
- Temporary space


- **Management**

- Space allocated when enter procedure
 - “Set-up” code
- Deallocated when return
 - “Finish” code

- **Pointers**


- Stack pointer `%esp` indicates stack top
- Frame pointer `%ebp` indicates start of current frame








Call Chain


yoo







Call Chain








Call Chain








Call Chain








Call Chain








Call Chain







Call Chain








Call Chain








Call Chain







Call Chain



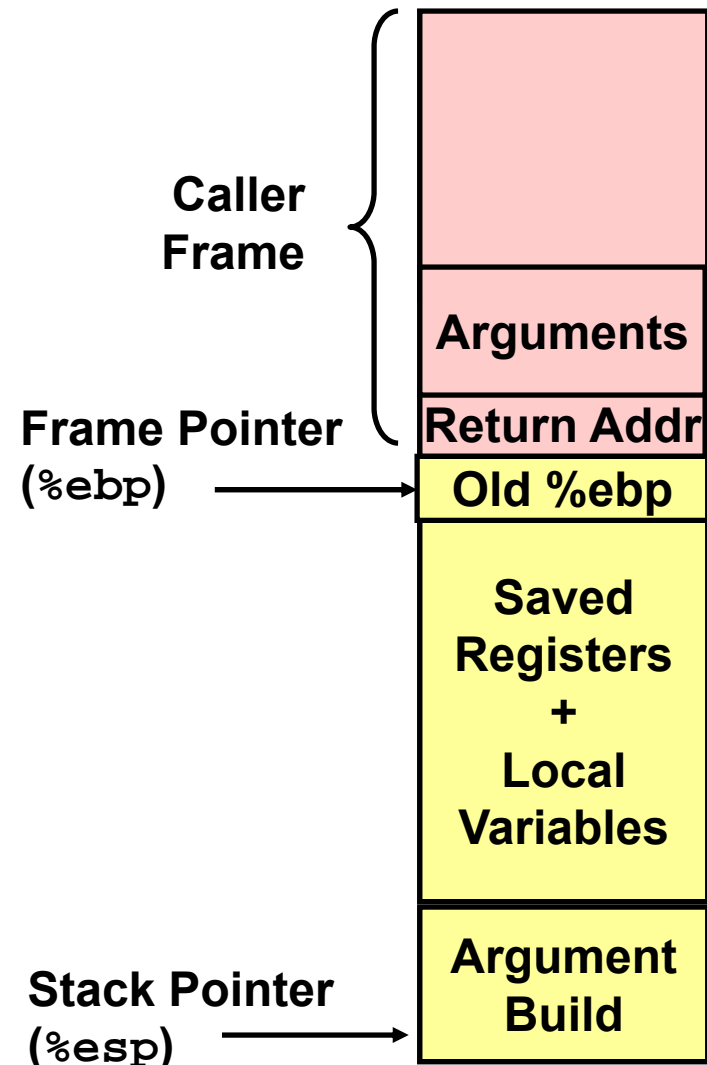


Call Chain



IA32/Linux Stack Frame

- Current Stack Frame (“Top” to Bottom)
 - Parameters for function about to call
 - “Argument build”
 - Local variables
 - If can’t keep in registers
 - Saved register context
 - Old frame pointer
- Caller Stack Frame
 - Return address
 - Pushed by `call` instruction
 - Arguments for this call



Revisiting swap

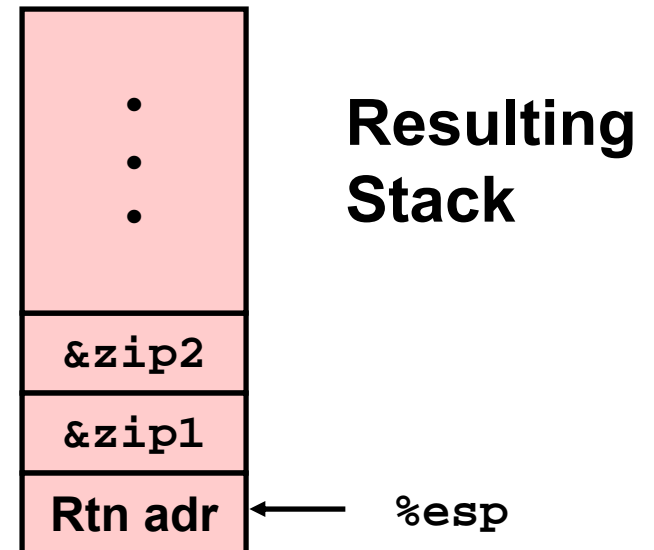
```
int zip1 = 15213;
int zip2 = 91125;

void call_swap()
{
    swap(&zip1, &zip2);
}
```

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

Calling swap from call_swap

```
call_swap:
    . . .
    pushl $zip2    # Global Var
    pushl $zip1    # Global Var
    call swap
    . . .
```



Revisiting swap

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

```
    pushl %ebp
    movl %esp,%ebp
    pushl %ebx
```

} Set Up

```
    movl 12(%ebp),%ecx
    movl 8(%ebp),%edx
    movl (%ecx),%eax
    movl (%edx),%ebx
    movl %eax, (%edx)
    movl %ebx, (%ecx)
```


} Body

```
    movl -4(%ebp),%ebx
    movl %ebp,%esp
    popl %ebp
    ret
```


} Finish

swap Setup #1

Entering Stack



Resulting Stack




swap:


```
pushl %ebp  
movl %esp,%ebp  
pushl %ebx
```

swap Setup #2

Entering Stack



Resulting Stack




swap:

```
    pushl %ebp  
    movl %esp,%ebp  
    pushl %ebx
```

swap Setup #3


Entering Stack



swap:


```
pushl %ebp  
movl %esp,%ebp  
pushl %ebx
```

Resulting Stack



Effect of swap Setup

Entering Stack



Offset
(relative to `%ebp`)


12

8

4

0

Resulting Stack

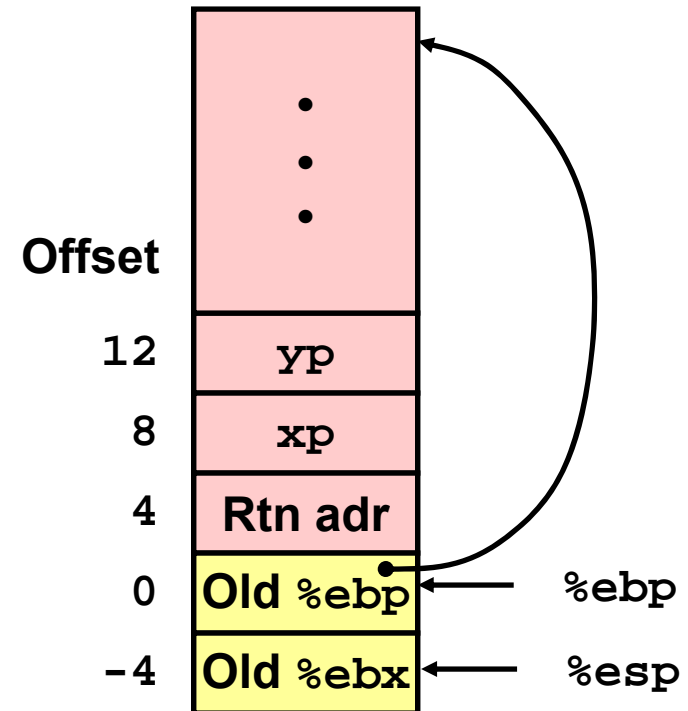
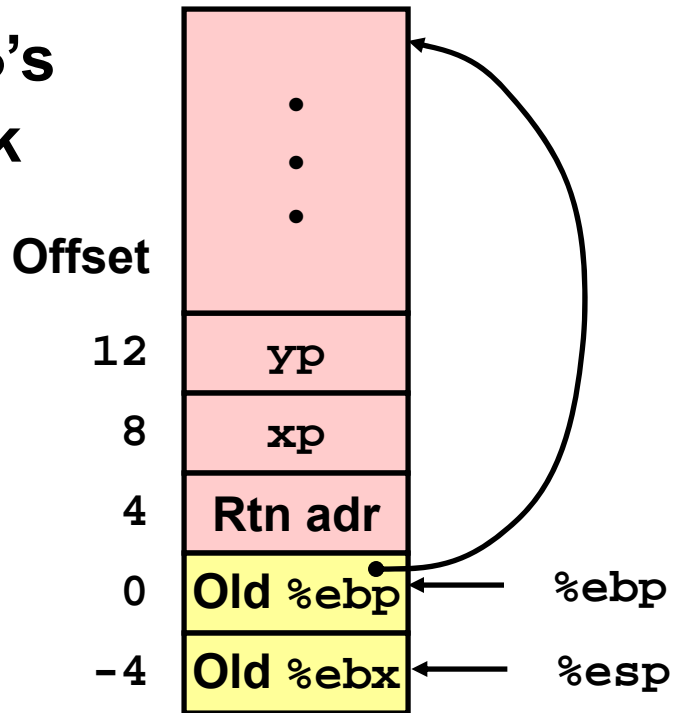


```
movl 12(%ebp),%ecx # get yp  
movl 8(%ebp),%edx  # get xp  
... }
```

Body

swap Finish #1

swap's
Stack




```
movl -4(%ebp), %ebx  
movl %ebp, %esp  
popl %ebp  
ret
```

- Observation


- Saved & restored register `%ebx`

swap Finish #2

swap's
Stack




swap's
Stack




```
movl -4(%ebp),%ebx  
movl %ebp,%esp  
popl %ebp  
ret
```

swap Finish #3

**swap's
Stack**

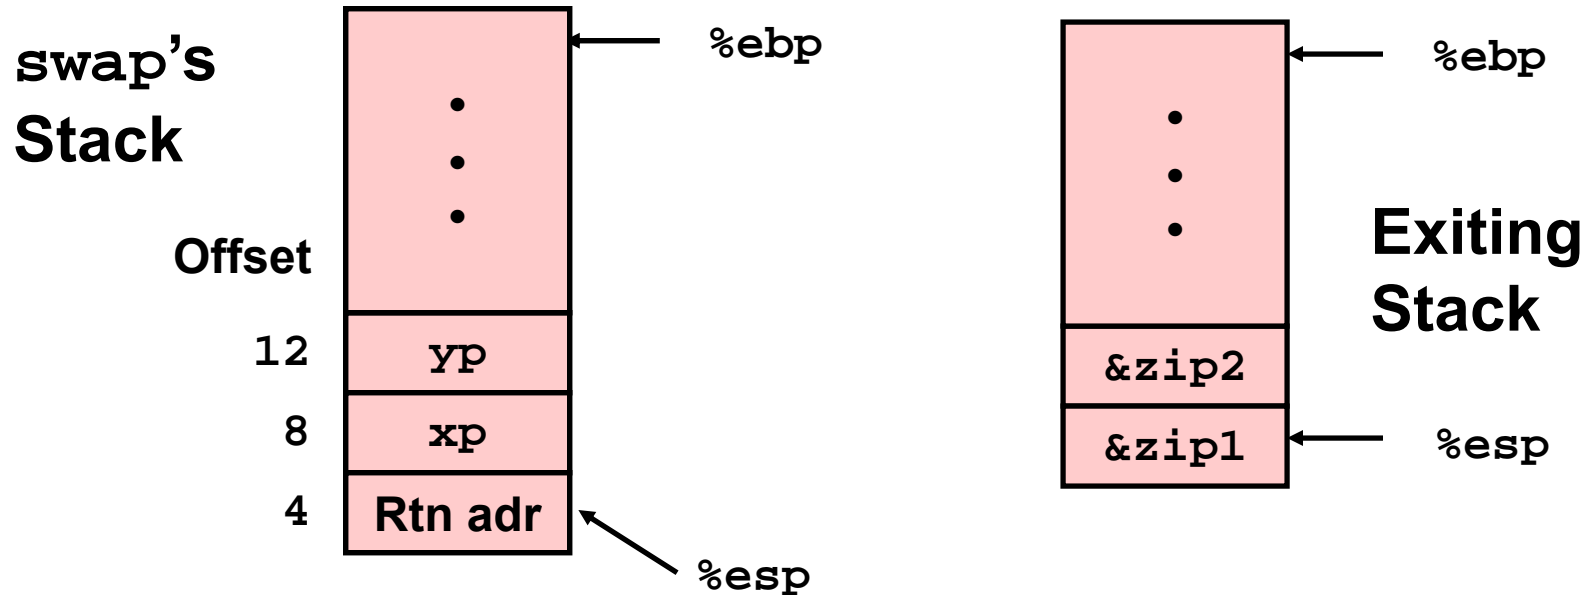


**swap's
Stack**



```
movl -4(%ebp),%ebx
movl %ebp,%esp
popl %ebp
ret
```

swap Finish #4



- **Observation**

- Saved & restored register %ebx
- Didn't do so for %eax, %ecx, or %edx

```
movl -4(%ebp),%ebx
movl %ebp,%esp
popl %ebp
ret
```

Register Saving Conventions

- When procedure *yoo* calls *who*:
 - *yoo* is the *caller*, *who* is the *callee*
- Can Registers be Used for Temporary Storage?

```
yoo:
    . . .
    movl $15213, %edx
    call who
    addl %edx, %eax
    . . .
    ret
```

```
who:
    . . .
    movl 8(%ebp), %edx
    addl $91125, %edx
    . . .
    ret
```

- Contents of register `%edx` overwritten by *who*

Register Saving Conventions

- When procedure *yoo* calls who:
 - *yoo* is the *caller*, who is the *callee*
- Can Registers be Used for Temporary Storage?
- Conventions
 - “Caller Save”
 - Caller saves temporary in its frame before calling
 - “Callee Save”
 - Callee saves temporary in its frame before using

IA32/Linux Register Usage

- Integer Registers

- Two have special uses

`%ebp, %esp`

- Three managed as callee-save

`%ebx, %esi, %edi`

- Old values saved on stack prior to using

- Three managed as caller-save

`%eax, %edx, %ecx`

- Do what you please, but expect any callee to do so, as well

- Register `%eax` also stores returned value

**Caller-Save
Temporaries**

**Callee-Save
Temporaries**

Special

`%eax`

`%edx`

`%ecx`

`%ebx`

`%esi`

`%edi`

`%esp`

`%ebp`

Recursive Factorial

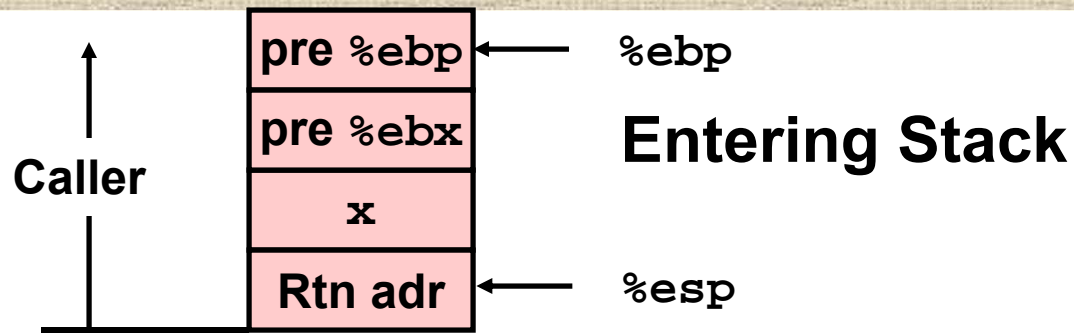
```
int rfact(int x)
{
    int rval;
    if (x <= 1)
        return 1;
    rval = rfact(x-1);
    return rval * x;
}
```

- Registers

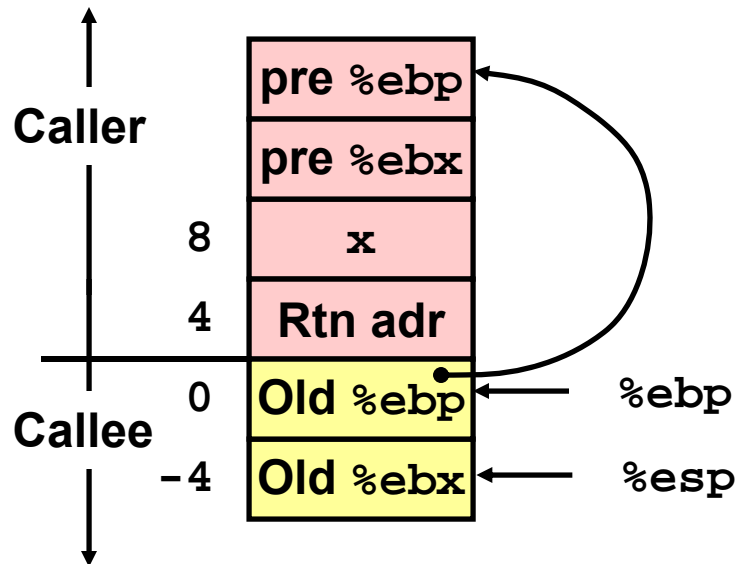
- %eax used without first saving
- %ebx used, but save at beginning & restore at end

```
.globl rfact
.type
rfact,@function
rfact:
    pushl %ebp
    movl %esp,%ebp
    pushl %ebx
    movl 8(%ebp),%ebx
    cmpl $1,%ebx
    jle .L78
    leal -1(%ebx),%eax
    pushl %eax
    call rfact
    imull %ebx,%eax
    jmp .L79
    .align 4
.L78:
    movl $1,%eax
.L79:
    movl -4(%ebp),%ebx
    movl %ebp,%esp
    popl %ebp
    ret
```


Rfact Stack Setup



```
rfact:
    pushl %ebp
    movl %esp,%ebp
    pushl %ebx
```



Rfact Body

Recursion

```
movl 8(%ebp),%ebx    # ebx = x
cmpl $1,%ebx         # Compare x : 1
jle .L78             # If <= goto Term
leal -1(%ebx),%eax    # eax = x-1
pushl %eax           # Push x-1
call rfact            # rfact(x-1)
imull %ebx,%eax       # rval * x
jmp .L79              # Goto done
.L78:                 # Term:
    movl $1,%eax      # return val = 1
.L79:                 # Done:
```

```
int rfact(int x)
{
    int rval;
    if (x <= 1)
        return 1;
    rval = rfact(x-1) ;
    return rval * x;
}
```

• Registers

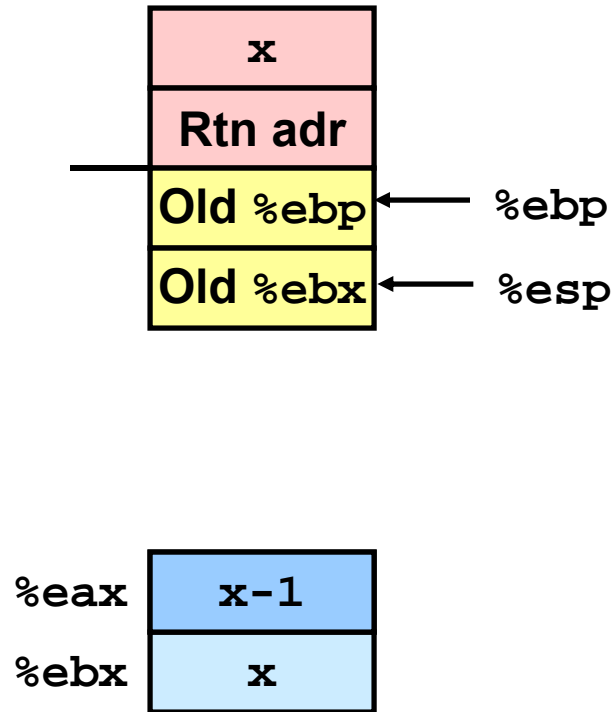
%ebx Stored value of x

%eax

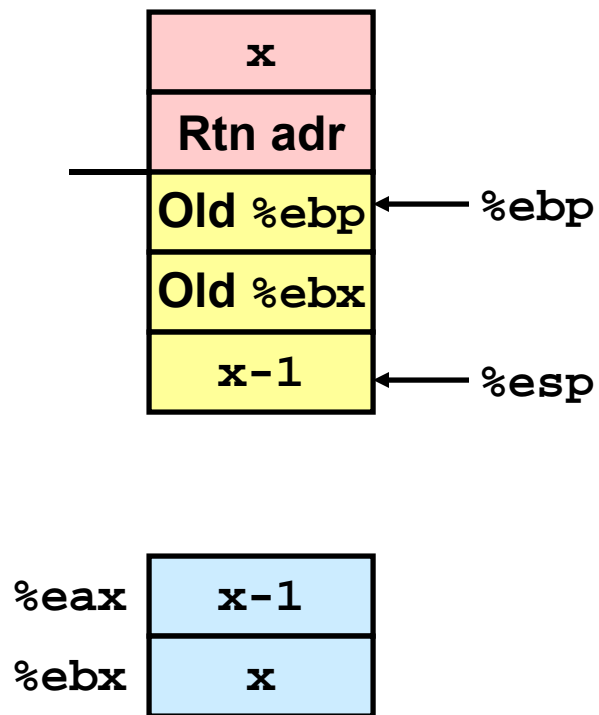
- Temporary value of x-1
- Returned value from rfact(x-1)
- Returned value from this call

Rfact Recursion

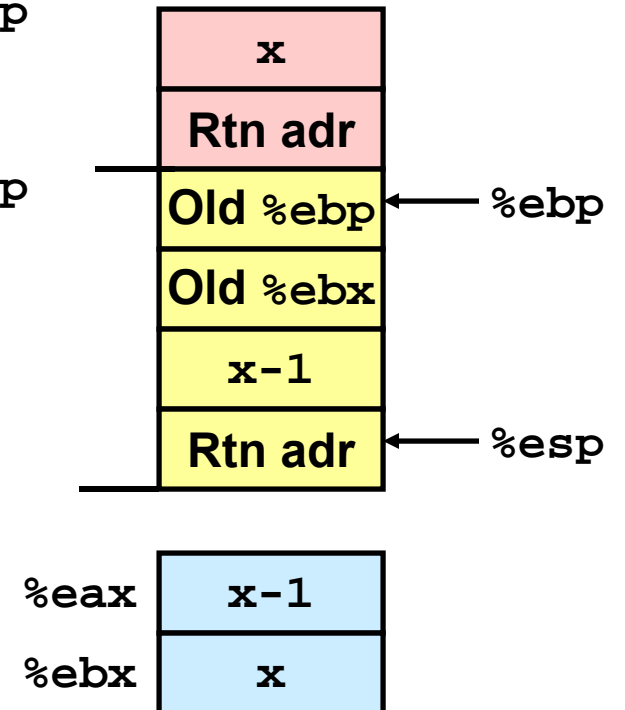
```
leal -1(%ebx),%eax
```



```
pushl %eax
```


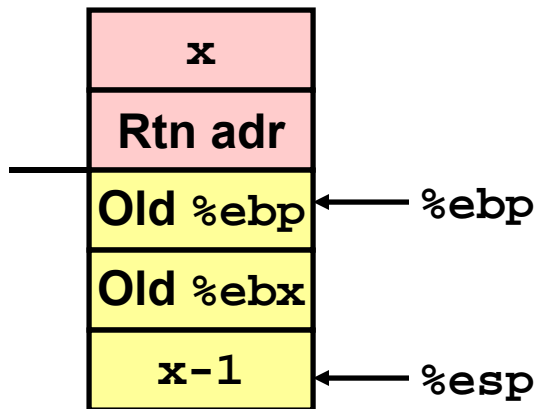


```
call rfact
```





Rfact Result

Return from Call

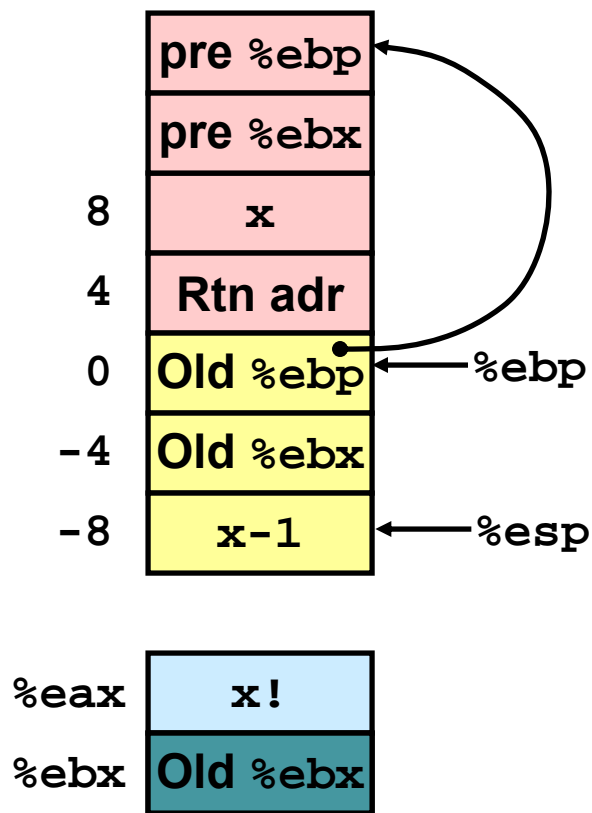


Assume that `rfact(x-1)` returns $(x-1)!$ in register %eax

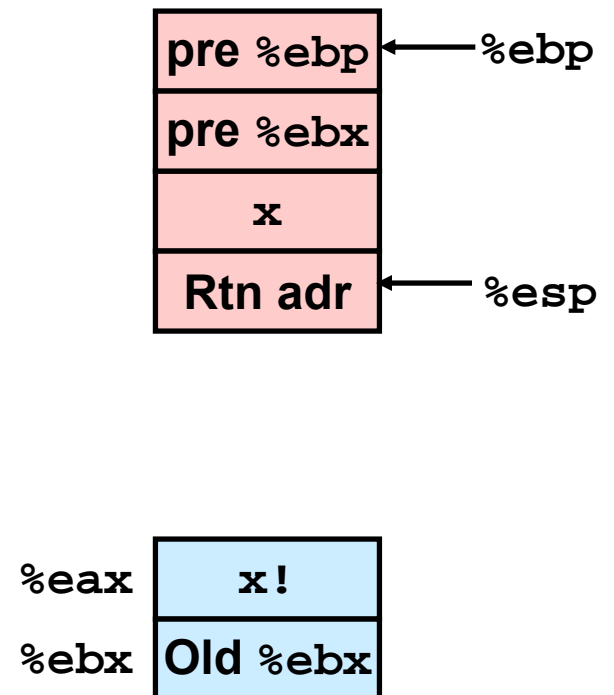
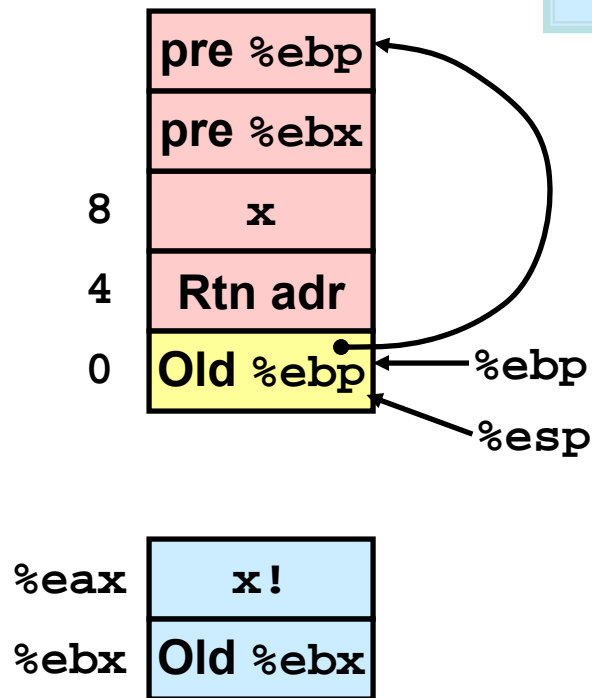
`imull %ebx,%eax`



Rfact Completion



```
movl -4(%ebp),%ebx  
movl %ebp,%esp  
popl %ebp  
ret
```



Pointer Code

Recursive Procedure

```
void s_helper
(int x, int *accum)
{
    if (x <= 1)
        return;
    else {
        int z = *accum * x;
        *accum = z;
        s_helper (x-1, accum);
    }
}
```

Top-Level Call

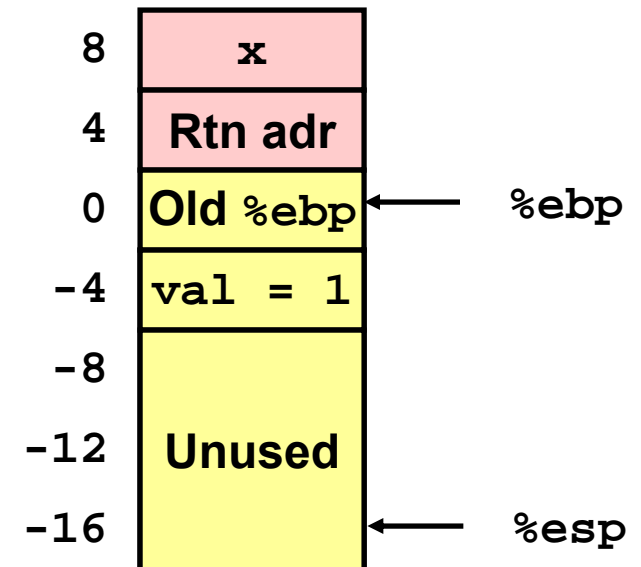
```
int sfact(int x)
{
    int val = 1;
    s_helper(x, &val);
    return val;
}
```

- Pass pointer to update location

Creating & Initializing Pointer

Initial part of sfact

```
_sfact:  
    pushl %ebp          # Save %ebp  
    movl %esp,%ebp      # Set %ebp  
    subl $16,%esp       # Add 16 bytes  
    movl 8(%ebp),%edx    # edx = x  
    movl $1,-4(%ebp)    # val = 1
```



- Using Stack for Local Variable

- Variable `val` must be stored on stack

- Need to create pointer to it

- Compute pointer as `-4(%ebp)`
- Push on stack as second argument

```
int sfact(int x)  
{  
    int val = 1;  
    s_helper(x, &val);  
    return val;  
}
```

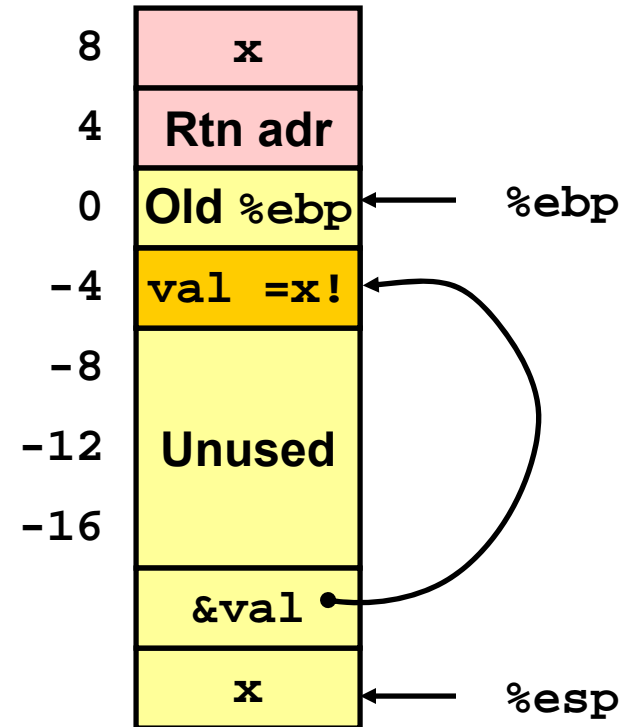

Passing Pointer

Calling s_helper from sfact

```
leal -4(%ebp),%eax # Compute &val
pushl %eax          # Push on stack
pushl %edx          # Push x
call s_helper       # call
movl -4(%ebp),%eax  # Return val
...                # Finish
```


```
int sfact(int x)
{
    int val = 1;
    s_helper(x, &val);
    return val;
}
```

Stack at time of call



Using Pointer

```
void s_helper
(int x, int *accum)
{
    . . .
    int z = *accum * x;
    *accum = z;
    . . .
}
```



```
. . .
movl %ecx,%eax      # z = x
imull (%edx),%eax   # z *= *accum
movl %eax,(%edx)    # *accum = z
. . .
```

- Register `%ecx` holds `x`
- Register `%edx` holds pointer to `accum`
 - Use access `(%edx)` to reference memory

Summary

- The Stack Makes Recursion Work
 - Private storage for each *instance* of procedure call
 - Instantiations don't clobber each other
 - Addressing of locals + arguments can be relative to stack positions
 - Can be managed by stack discipline
 - Procedures return in inverse order of calls
- IA32 Procedures Combination of Instructions + Conventions
 - Call / Ret instructions
 - Register usage conventions
 - Caller / Callee save
 - %ebp and %esp
 - Stack frame organization conventions

x86-64 Integer Registers

%rax	%eax
%rbx	%ebx
%rcx	%ecx
%rdx	%edx
%rsi	%esi
%rdi	%edi
%rsp	%esp
%rbp	%ebp

%r8	%r8d
%r9	%r9d
%r10	%r10d
%r11	%r11d
%r12	%r12d
%r13	%r13d
%r14	%r14d
%r15	%r15d

- Twice the number of registers
- Accessible as 8, 16, 32, 64 bits

x86-64 Integer Registers

%rax	Return value
%rbx	Callee saved
%rcx	Argument #4
%rdx	Argument #3
%rsi	Argument #2
%rdi	Argument #1
%rsp	Stack pointer
%rbp	Callee saved

%r8	Argument #5
%r9	Argument #6
%r10	Callee saved
%r11	Used for linking
%r12	C: Callee saved
%r13	Callee saved
%r14	Callee saved
%r15	Callee saved

x86-64 Registers

- Arguments passed to functions via registers
 - If more than 6 integral parameters, then pass rest on stack
 - These registers can be used as caller-saved as well
- All references to stack frame via stack pointer
 - Eliminates need to update `%ebp/%rbp`
- Other Registers
 - 6+1 callee saved
 - 2 or 3 have special uses

x86-64 Long Swap

```
void swap(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

```
swap:
    movq    (%rdi), %rdx
    movq    (%rsi), %rax
    movq    %rax, (%rdi)
    movq    %rdx, (%rsi)
    ret
```

- Operands passed in registers
 - First (**x**p) in %**rdi**, second (**y**p) in %**rsi**
 - 64-bit pointers
- No stack operations required (except `ret`)
- Avoiding stack
 - Can hold all local information in registers


x86-64 Locals in the Red Zone

```
/* Swap, using local array */
void swap_a(long *xp, long *yp)
{
    volatile long loc[2];
    loc[0] = *xp;
    loc[1] = *yp;
    *xp = loc[1];
    *yp = loc[0];
}
```

swap_a:

```
movq    (%rdi), %rax
movq    %rax, -24(%rsp)
movq    (%rsi), %rax
movq    %rax, -16(%rsp)
movq    -16(%rsp), %rax
movq    %rax, (%rdi)
movq    -24(%rsp), %rax
movq    %rax, (%rsi)
ret
```

- Avoiding Stack Pointer Change
 - Can hold all information within small window beyond stack pointer



x86-64 NonLeaf without Stack Frame

```
long scount = 0;

/* Swap a[i] & a[i+1] */
void swap_ele_se
    (long a[], int i)
{
    swap(&a[i], &a[i+1]);
    scount++;
}
```

- No values held while swap being invoked
- No callee save registers needed

```
swap_ele_se:
    movslq %esi,%rsi          # Sign extend i
    leaq    (%rdi,%rsi,8), %rdi # &a[i]
    leaq    8(%rdi), %rsi      # &a[i+1]
    call    swap              # swap()
    incq    scount(%rip)       # scount++;
    ret
```


x86-64 Call using Jump

```
long scount = 0;

/* Swap a[i] & a[i+1] */
void swap_ele(long a[], int i)
{
    swap(&a[i], &a[i+1]);
}
```

- When swap executes ret, it will return from swap_ele
- Possible since swap is a “tail call”
(no instructions afterwards)

swap_ele:

```
    movslq %esi,%rsi          # Sign extend i
    leaq    (%rdi,%rsi,8), %rdi # &a[i]
    leaq    8(%rdi), %rsi      # &a[i+1]
    jmp     swap               # swap()
```

x86-64 Stack Frame Example

```
long sum = 0;
/* Swap a[i] & a[i+1] */
void swap_ele_su
    (long a[], int i)
{
    swap(&a[i], &a[i+1]);
    sum += a[i];
}
```

- Keeps values of `a` and `i` in callee save registers
- Must set up stack frame to save these registers

`swap_ele_su:`

```
movq    %rbx, -16(%rsp)
movslq  %esi,%rbx
movq    %r12, -8(%rsp)
movq    %rdi, %r12
leaq    (%rdi,%rbx,8), %rdi
subq    $16, %rsp
leaq    8(%rdi), %rsi
call   swap
movq    (%r12,%rbx,8), %rax
addq    %rax, sum(%rip)
movq    (%rsp), %rbx
movq    8(%rsp), %r12
addq    $16, %rsp
ret
```

Understanding x86-64 Stack Frame

swap_ele_su:

<u>movq</u>	<u>%rbx, -16(%rsp)</u>	# Save %rbx
movslq	%esi,%rbx	# Extend & save i
<u>movq</u>	<u>%r12, -8(%rsp)</u>	# Save %r12
movq	%rdi, %r12	# Save a
leaq	(%rdi,%rbx,8), %rdi	# &a[i]
<u>subq</u>	<u>\$16, %rsp</u>	# Allocate stack frame
leaq	8(%rdi), %rsi	# &a[i+1]
call	swap	# swap()
movq	(%r12,%rbx,8), %rax	# a[i]
addq	%rax, sum(%rip)	# sum += a[i]
<u>movq</u>	<u>(%rsp), %rbx</u>	# Restore %rbx
<u>movq</u>	<u>8(%rsp), %r12</u>	# Restore %r12
<u>addq</u>	<u>\$16, %rsp</u>	# Deallocate stack frame
ret		

Understanding x86-64 Stack Frame

swap_ele_su:

movq %rbx, -16(%rsp)

Save %rbx

movq %r12, -8(%rsp)

Save %r12

subq \$16, %rsp

Allocate stack frame

movq (%rsp), %rbx



Restore %rbx

movq 8(%rsp), %r12

Restore %r12

addq \$16, %rsp

Deallocate stack frame



Interesting Features of Stack Frame

- Allocate entire frame at once
 - All stack accesses can be relative to `%rsp`
 - Do by decrementing stack pointer
 - Can delay allocation, since safe to temporarily use red zone
- Simple deallocation
 - Increment stack pointer
 - No base/frame pointer needed

x86-64 Procedure Summary

- Heavy use of registers
 - Parameter passing
 - More temporaries since more registers
- Minimal use of stack
 - Sometimes none
 - Allocate/deallocate entire block
- Many tricky optimizations
 - What kind of stack frame to use
 - Calling with jump
 - Various allocation techniques