

# Machine-Level Programming: Miscellaneous Topics


## Chapter 3 of B&O

Some notes adopted from Bryant and O'Hallaron

# IA32 Linux Memory Layout

- **Stack**
  - Runtime stack (8MB limit)
- **Heap**
  - Dynamically allocated storage
  - When call `malloc()` , `calloc()` , `new()`
- **Data**
  - Statically allocated data
  - E.g., arrays & strings declared in code
- **Text**
  - Executable machine instructions
  - Read-only

Upper 2 hex digits  
= 8 bits of address



# Memory Allocation Example

```
char big_array[1<<24]; /* 16 MB */
char huge_array[1<<28]; /* 256 MB */


int beyond;
char *p1, *p2, *p3, *p4;

int useless() { return 0; }

int main()
{
    p1 = malloc(1 << 28); /* 256 MB */
    p2 = malloc(1 << 8); /* 256 B */
    p3 = malloc(1 << 28); /* 256 MB */
    p4 = malloc(1 << 8); /* 256 B */
    /* Some print statements ... */
}
```

*Where does everything go?*

The pointers point to stuff in the heap, but they themselves are in data






# IA32 Example Addresses

*address range  $\sim 2^{32}$*

\$esp	0xffffbcd0
p3	0x65586008
p1	0x55585008
p4	0x1904a110
p2	0x1904a008
&p2	0x18049760
beyond	0x08049744
big_array	0x18049780
huge_array	0x08049760
main()	0x080483c6
useless()	0x08049744
final malloc()	0x006be166

malloc() is dynamically linked  
address determined at runtime




# x86-64 Example Addresses

address range  $\sim 2^{47}$

<code>\$rsp</code>	<code>0x7fffffff8d1f8</code>
<code>p3</code>	<code>0x2aaabaadd010</code>
<code>p1</code>	<code>0x2aaaaadc010</code>
<code>p4</code>	<code>0x000011501120</code>
<code>p2</code>	<code>0x000011501010</code>
<code>&amp;p2</code>	<code>0x000010500a60</code>
<code>beyond</code>	<code>0x000000500a44</code>
<code>big_array</code>	<code>0x000010500a80</code>
<code>huge_array</code>	<code>0x000000500a50</code>
<code>main()</code>	<code>0x000000400510</code>
<code>useless()</code>	<code>0x000000400500</code>
<code>final malloc()</code>	<code>0x00386ae6a170</code>

`malloc()` is dynamically linked  
address determined at runtime



# C operators

## Operators

`() [] -> .`  
`! ~ ++ -- + - * & (type) sizeof`  
`* / %`  
`+ -`  
`<< >>`  
`< <= > >=`  
`== !=`  
`&`  
`^`  
`|`  
`&&`  
`||`  
`?:`  
`= += -= *= /= %= &= ^= != <<= >>=`  
`,`

## Associativity

left to right  
right to left  
left to right  
left to right  
left to right  
left to right  
left to right  
left to right  
left to right  
left to right  
right to left  
right to left  
left to right

**Note: Unary +, -, and \* have higher precedence than binary forms**

# C pointer declarations

```
int *p
```

**p is a pointer to int**

```
int *p[13]
```

**p is an array[13] of pointer to int**

```
int *(p[13])
```

**p is an array[13] of pointer to int**

```
int **p
```

**p is a pointer to a pointer to an int**

```
int (*p)[13]
```

**p is a pointer to an array[13] of int**

```
int *f()
```

**f is a function returning a pointer to int**

```
int (*f)()
```

**f is a pointer to a function returning int**

```
int ((*f())[13])()
```

**f is a function returning ptr to an array[13] of pointers to functions returning int**

```
int ((*x[3])())[5]
```

**x is an array[3] of pointers to functions returning pointers to array[5] of ints**


# Internet Worm and IM War

- November, 1988

- Internet Worm attacks thousands of Internet hosts.
- How did it happen?

- July, 1999

- Microsoft launches MSN Messenger (instant messaging system).
- Messenger clients can access popular AOL Instant Messaging Service (AIM) servers





# Internet Worm and IM War (cont.)

- August 1999
  - Mysteriously, Messenger clients can no longer access AIM servers.
  - Microsoft and AOL begin the IM war:
    - AOL changes server to disallow Messenger clients
    - Microsoft makes changes to clients to defeat AOL changes.
    - At least 13 such skirmishes.
  - How did it happen?
- The Internet Worm and AOL/Microsoft War were both based on *stack buffer overflow* exploits!
  - many Unix functions do not check argument sizes.
  - allows target buffers to overflow.

# String Library Code

- Implementation of Unix function `gets`

- No way to specify limit on number of characters to read

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getc();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getc();
    }
    *p = '\0';
    return dest;
}
```

- Similar problems with other Unix functions

- `strcpy`: Copies string of arbitrary length
- `scanf`, `fscanf`, `sscanf`,
  - when given `%s` conversion specification

# Vulnerable Buffer Code

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
int main()  
{  
    printf("Type a string:");  
    echo();  
    return 0;  
}
```

# Buffer Overflow Executions

```
unix>./bufdemo  
Type a string:123  
123
```


```
unix>./bufdemo  
Type a string:12345  
Segmentation Fault
```

```
unix>./bufdemo  
Type a string:12345678  
Segmentation Fault
```



# Buffer Overflow Stack

*Before call to gets*



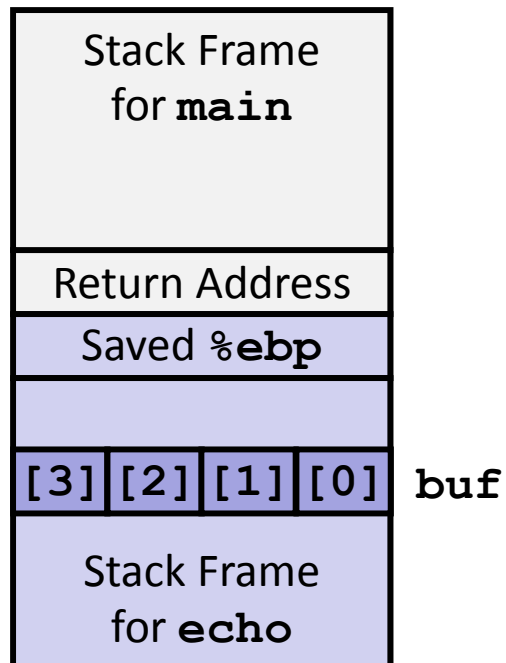
```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

```
echo:
    pushl %ebp                # Save %ebp on stack
    movl  %esp, %ebp
    pushl %ebx                # Save %ebx
    leal  -8(%ebp), %ebx      # Compute buf as %ebp-8
    subl  $20, %esp           # Allocate stack space
    movl  %ebx, (%esp)        # Push buf on stack
    call  gets                # Call gets
    . . .
```


# Buffer Overflow Stack Example

```
unix> gdb bufdemo
(gdb) break echo
Breakpoint 1 at 0x8048583
(gdb) run
Breakpoint 1, 0x8048583 in echo ()
(gdb) print /x $ebp
$1 = 0xffffc638
(gdb) print /x *(unsigned *)$ebp
$2 = 0xffffc658
(gdb) print /x *((unsigned *)$ebp + 1)
$3 = 0x80485f7
```

*Before call to gets*



*Before call to gets*




80485f2: call 80484f0 <echo>


80485f7: mov 0xffffffffc(%ebp), %ebx # Return Point

# Buffer Overflow Example #1

*Before call to gets*




*Input 1234567*



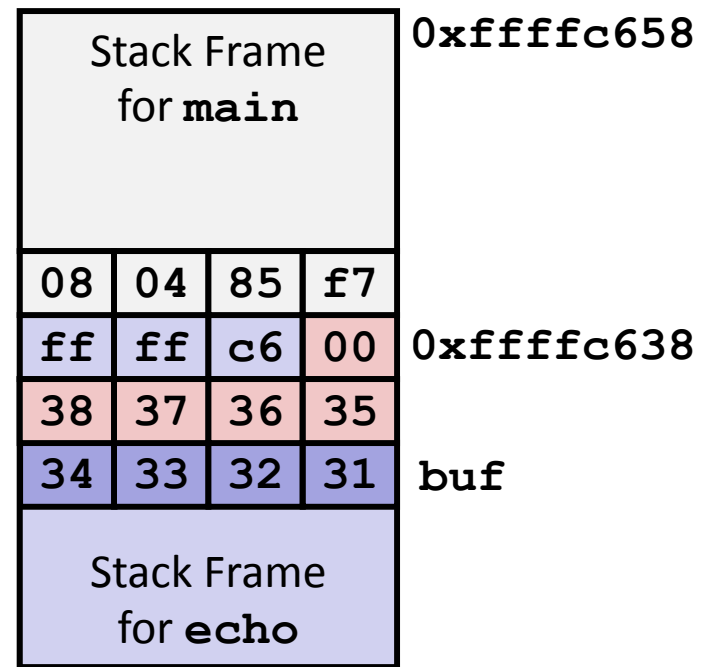
**Overflow buf, but no problem**

# Buffer Overflow Example #2

*Before call to gets*



*Input 12345678*




**Base pointer corrupted**

```
. . .
804850a: 83 c4 14  add    $0x14,%esp  # deallocate space
804850d: 5b        pop     %ebx      # restore %ebx
804850e: c9        leave   # movl %ebp, %esp; popl %ebp
804850f: c3        ret      # Return
```




# Buffer Overflow Example #3

*Before call to gets*



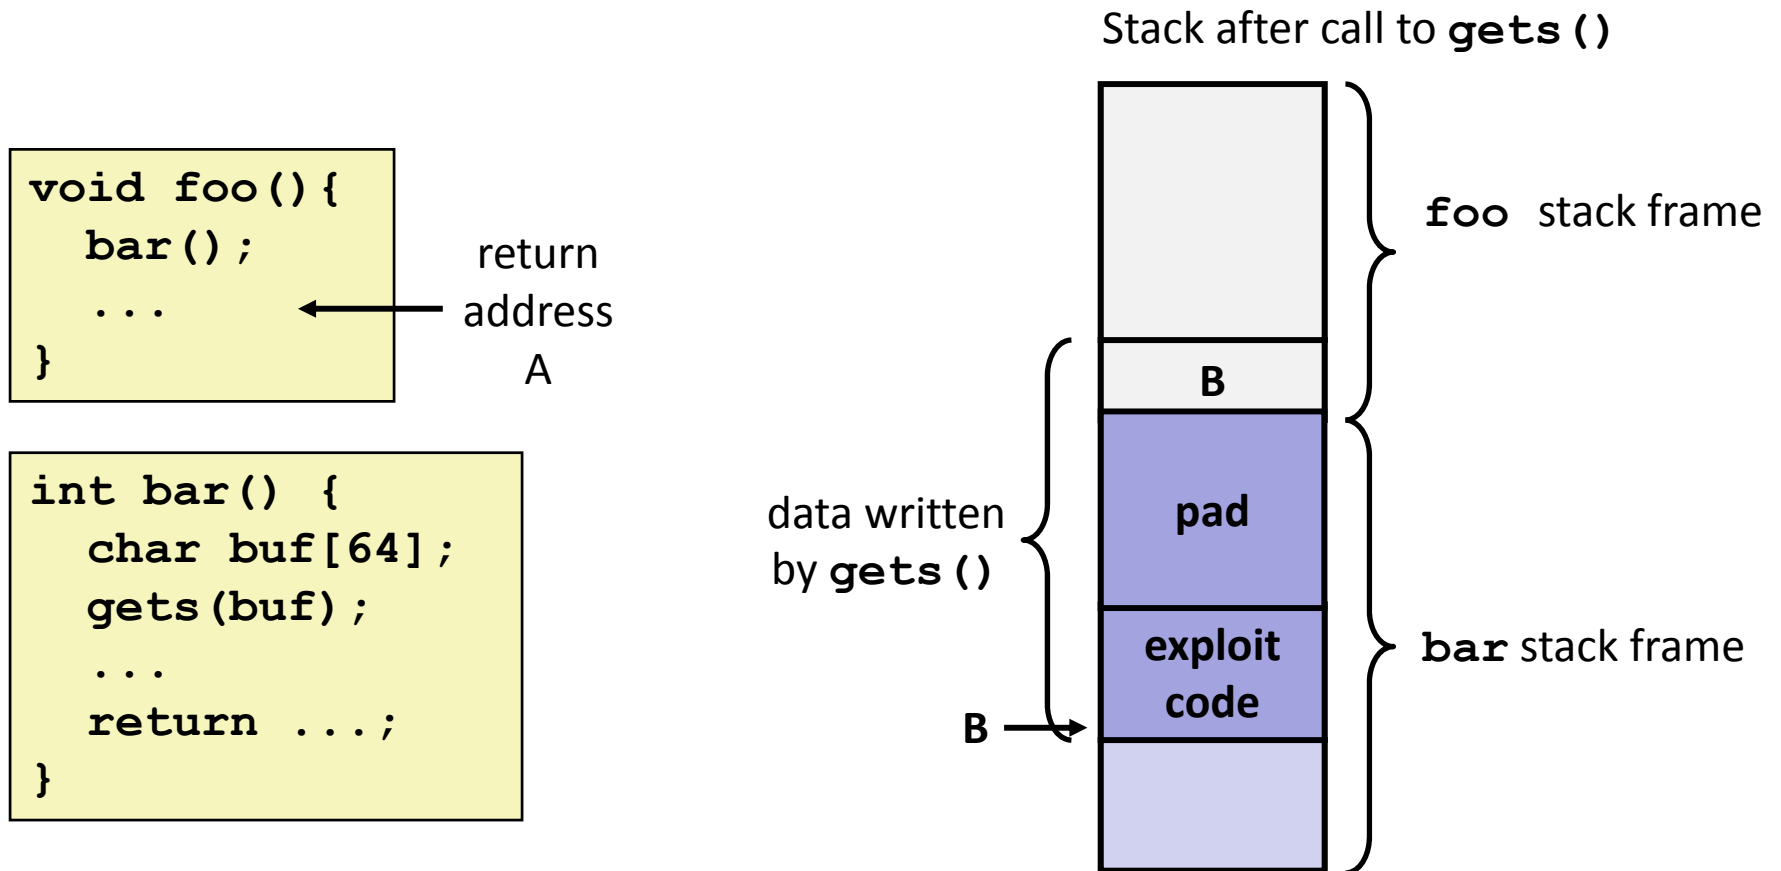
*Input 12345678*



**Return address corrupted**

```
80485f2: call 80484f0 <echo>
80485f7: mov 0xffffffffc(%ebp),%ebx # Return Point
```

# Malicious Use of Buffer Overflow



# Exploits Based on Buffer Overflows

- *Buffer overflow bugs allow remote machines to execute arbitrary code on victim machines.*
- Internet worm
  - Early versions of the finger server (fingerd) used `gets ( )` to read the argument sent by the client:
    - *finger droh@cs.cmu.edu*
  - Worm attacked fingerd server by sending phony argument:
    - *finger "exploit-code padding new-return-address"*
    - exploit code: executed a root shell on the victim machine with a direct TCP connection to the attacker.

# Exploits Based on Buffer Overflows

- *Buffer overflow bugs allow remote machines to execute arbitrary code on victim machines.*
- IM War
  - AOL exploited existing buffer overflow bug in AIM clients
  - exploit code: returned 4-byte signature (the bytes at some location in the AIM client) to server.
  - When Microsoft changed code to match signature, AOL changed signature location.



Date: Wed, 11 Aug 1999 11:30:57 -0700 (PDT)  
From: Phil Bucking <philbucking@yahoo.com>  
Subject: AOL exploiting buffer overrun bug in their own software!  
To: rms@pharlap.com

Mr. Smith,

I am writing you because I have discovered something that I think you might find interesting because you are an Internet security expert with experience in this area. I have also tried to contact AOL but received no response.

I am a developer who has been working on a revolutionary new instant messaging client that should be released later this year.

...

It appears that the AIM client has a buffer overrun bug. By itself this might not be the end of the world, as MS surely has had its share. But AOL is now \*exploiting their own buffer overrun bug\* to help in its efforts to block MS Instant Messenger.

....

Since you have significant credibility with the press I hope that you can use this information to help inform people that behind AOL's friendly exterior they are nefariously compromising peoples' security.

Sincerely,  
Phil Bucking  
Founder, Bucking Consulting  
philbucking@yahoo.com

**It was later determined that this email  
originated from within Microsoft!**

# Avoiding Overflow Vulnerability

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    fgets(buf, 4, stdin);  
    puts(buf);  
}
```

- Use Library Routines that Limit String Lengths
  - fgets instead of gets
  - strncpy instead of strcpy
  - Don't use scanf with %s conversion specification
    - Use fgets to read the string

# System-Level Protections

- Randomized stack offsets

- At start of program, allocate random amount of space on stack
- Makes it difficult for hacker to predict beginning of inserted code

- Nonexecutable code segments

- In traditional x86, can mark region of memory as either “read-only” or “writeable”
  - Can execute anything readable
- Add explicit “execute” permission

```
unix> gdb bufdemo
(gdb) break echo

(gdb) run
(gdb) print /x $ebp
$1 = 0xfffffc638

(gdb) run
(gdb) print /x $ebp
$2 = 0xffffbb08

(gdb) run
(gdb) print /x $ebp
$3 = 0xffffc6a8
```

# Final Observations

- Memory Layout
  - OS/machine dependent (including kernel version)
  - Basic partitioning: stack/data/text/heap/DLL found in most machines
- Type Declarations in C
  - Notation obscure, but very systematic
- Working with Strange Code
  - Important to analyze nonstandard cases
    - E.g., what happens when stack corrupted due to buffer overflow
  - Helps to step through with GDB