

Some OS Basics

Some notes adopted from Bryant and O'Hallaron

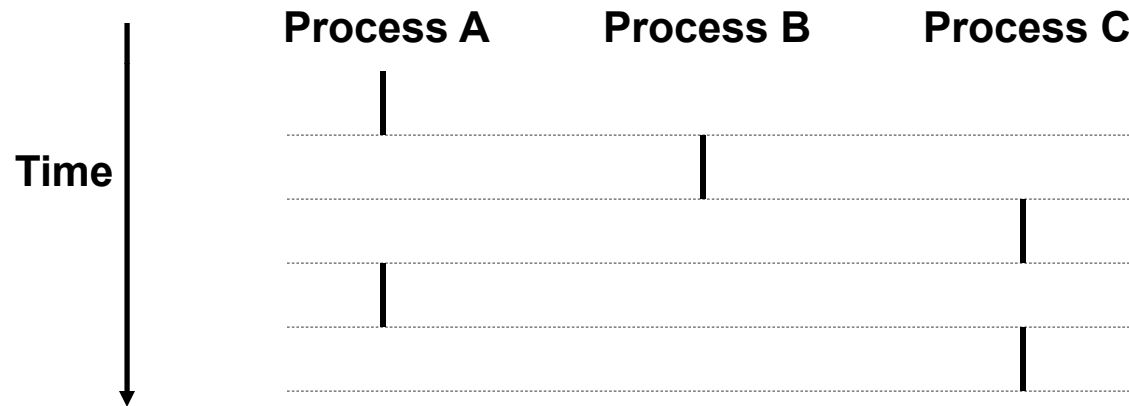
Role of The Operating System?

Processes

- Def: A process is an instance of a running program.
 - One of the most profound ideas in computer science.
 - Not the same as “program” or “processor”
- Process provides each program with two key abstractions:
 - Logical control flow
 - Each program seems to have exclusive use of the CPU.
 - Private address space
 - Each program seems to have exclusive use of main memory.
- How are these Illusions maintained?
 - Process executions interleaved (multitasking)
 - Address spaces managed by virtual memory system

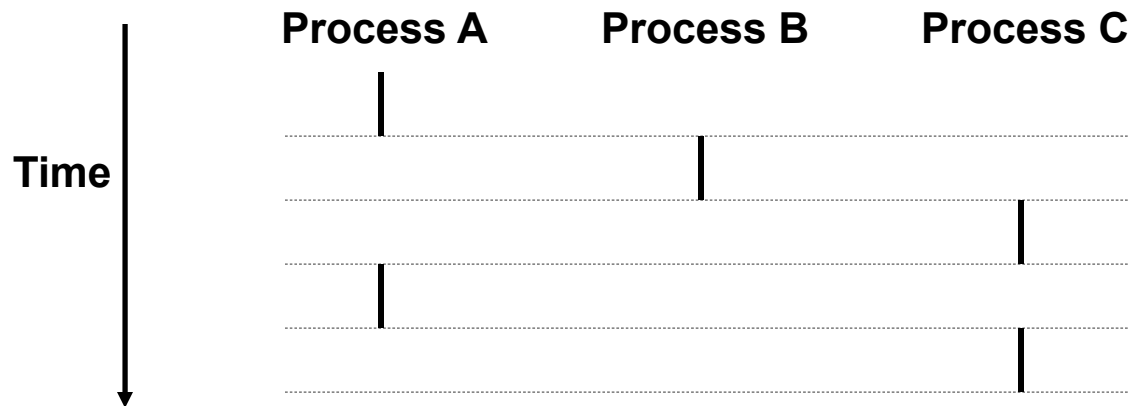
Logical Control Flows

Each process has its own logical control flow




Concurrent Processes

- Two processes *run concurrently (are concurrent)* if their flows overlap in time.
- Otherwise, they are *sequential*.
- Examples:
 - Concurrent: A & B, A & C
 - Sequential: B & C



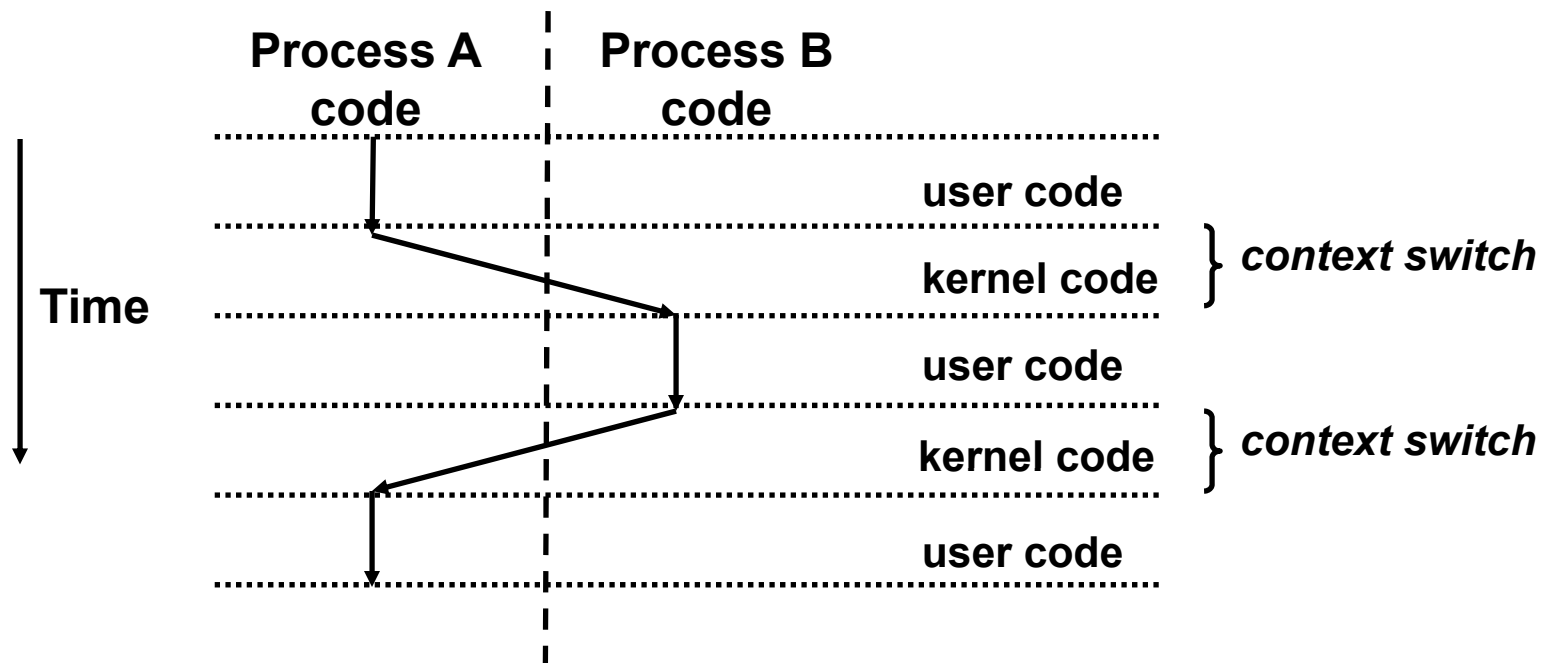
User View of Concurrent Processes

- Control flows for concurrent processes are physically disjoint in time.
- However, we can think of concurrent processes as running in parallel with each other.



Context Switching

- Processes are managed by a shared chunk of OS code called the *kernel*
- Control flow passes from one process to another via a *context switch*.



Process: Traditional View

- Process = process context + code, data, and stack

Process context


Program context:

Data registers
Condition codes
Stack pointer (SP)
Program counter (PC)

Kernel context:

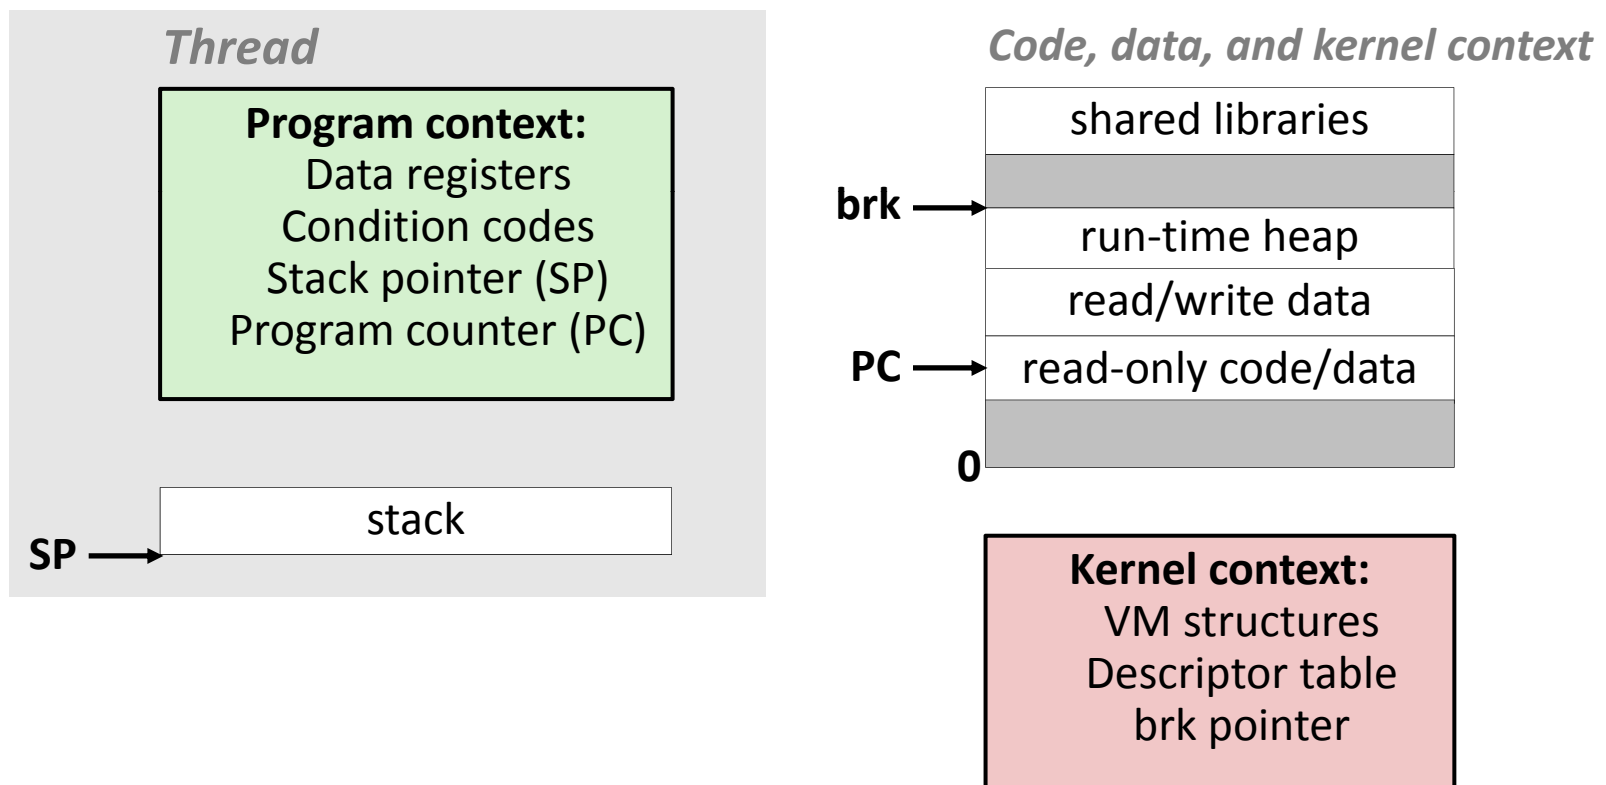
VM structures
Descriptor table
brk pointer

Code, data, and stack



Process: Alternative View

- Process = thread + code, data, and kernel context



Process with Two Threads

Thread 1

Program context:

Data registers
Condition codes
Stack pointer (SP)
Program counter (PC)




Thread 2

Program context:

Data registers
Condition codes
Stack pointer (SP)
Program counter (PC)



Code, data, and kernel context



Kernel context:

VM structures
Descriptor table
brk pointer


Threads vs. Processes

- Threads and processes: similarities
 - Each has its own logical control flow
 - Each can run concurrently with others
 - Each is context switched (scheduled) by the kernel
- Threads and processes: differences
 - Threads share code and data, processes (typically) do not
 - Threads are much less expensive than processes
 - Process control (creating and reaping) is more expensive as thread control
 - Context switches for processes much more expensive than for threads


Threads vs. Processes (contd.)

- Processes form a tree hierarchy
- Threads form a pool of peers
 - Each thread can kill any other
 - Each thread can wait for any other thread to terminate
 - Main thread: first thread to run in a process

Process hierarchy




Thread pool



Virtual Memory (Previous Lectures)


- Programs refer to virtual memory addresses
 - `movl (%ecx),%eax`
 - Conceptually very large array of bytes
 - Each byte has its own address
 - Actually implemented with hierarchy of different memory
 - System provides address space private to particular “process”
- Allocation: Compiler and run-time system
 - Where different program objects should be stored
 - All allocation within single virtual address space
- But why virtual memory?
- Why not physical memory?



Problem 1: How Does Everything Fit?

**64-bit addresses:
16 Exabyte**

**Physical main memory:
Few Gigabytes**



And there are many processes

Problem 2: Memory Management

Process 1
Process 2
Process 3
...
Process n

X


stack
heap
.text
.data
...

*What goes
where?*


Physical main memory



Problem 3: How To Protect




Problem 4: How To Share?



Solution: Level Of Indirection

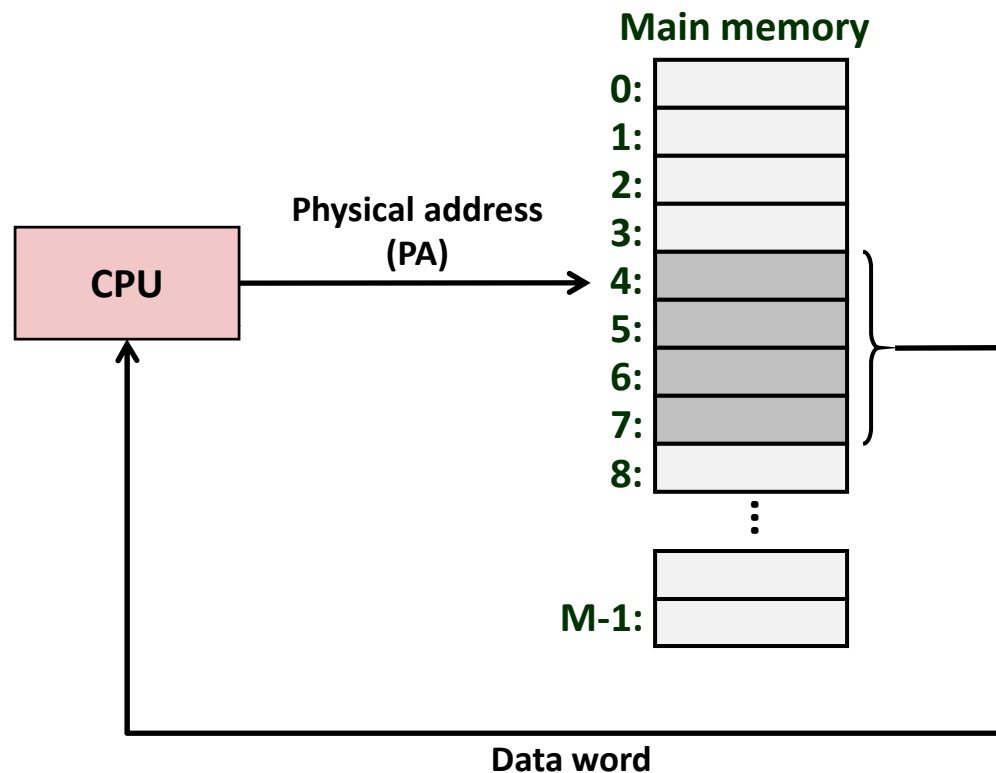
- Each process gets its own private memory space
- Solves the previous problems



Address Spaces

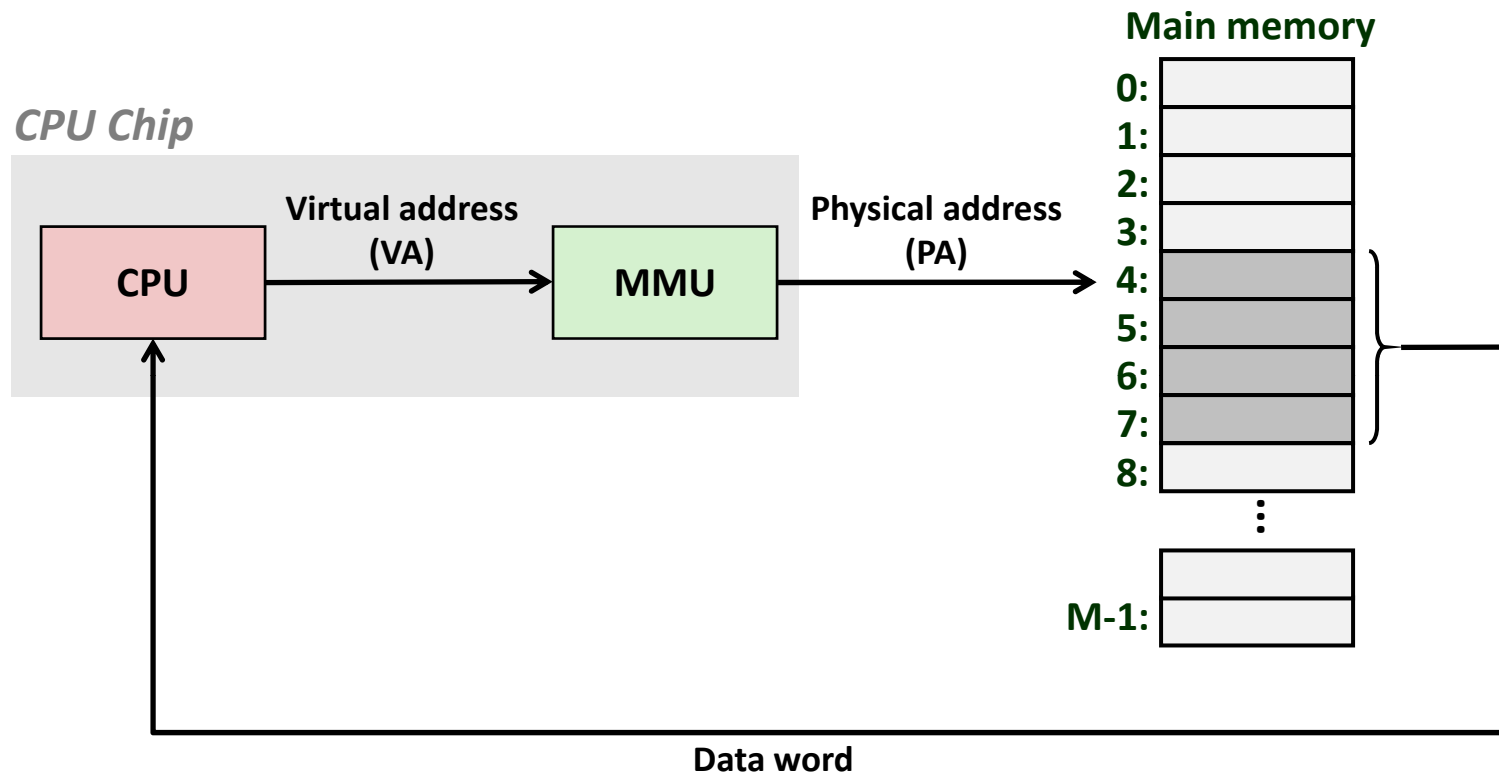
- Virtual address space
 - Set of $N = 2^n$ virtual addresses: $\{0, 1, 2, 3, \dots, N-1\}$
- Physical address space
 - Set of $M = 2^m$ physical addresses: $\{0, 1, 2, 3, \dots, M-1\}$
- Clean distinction between data (bytes) and their attributes (addresses)
- Each object can now have multiple addresses
- Every byte in main memory:
one physical address, one (or more) virtual addresses

A System Using Physical Addressing



- Used in “simple” systems like embedded microcontrollers in devices like cars, elevators, and digital picture frames

A System Using Virtual Addressing



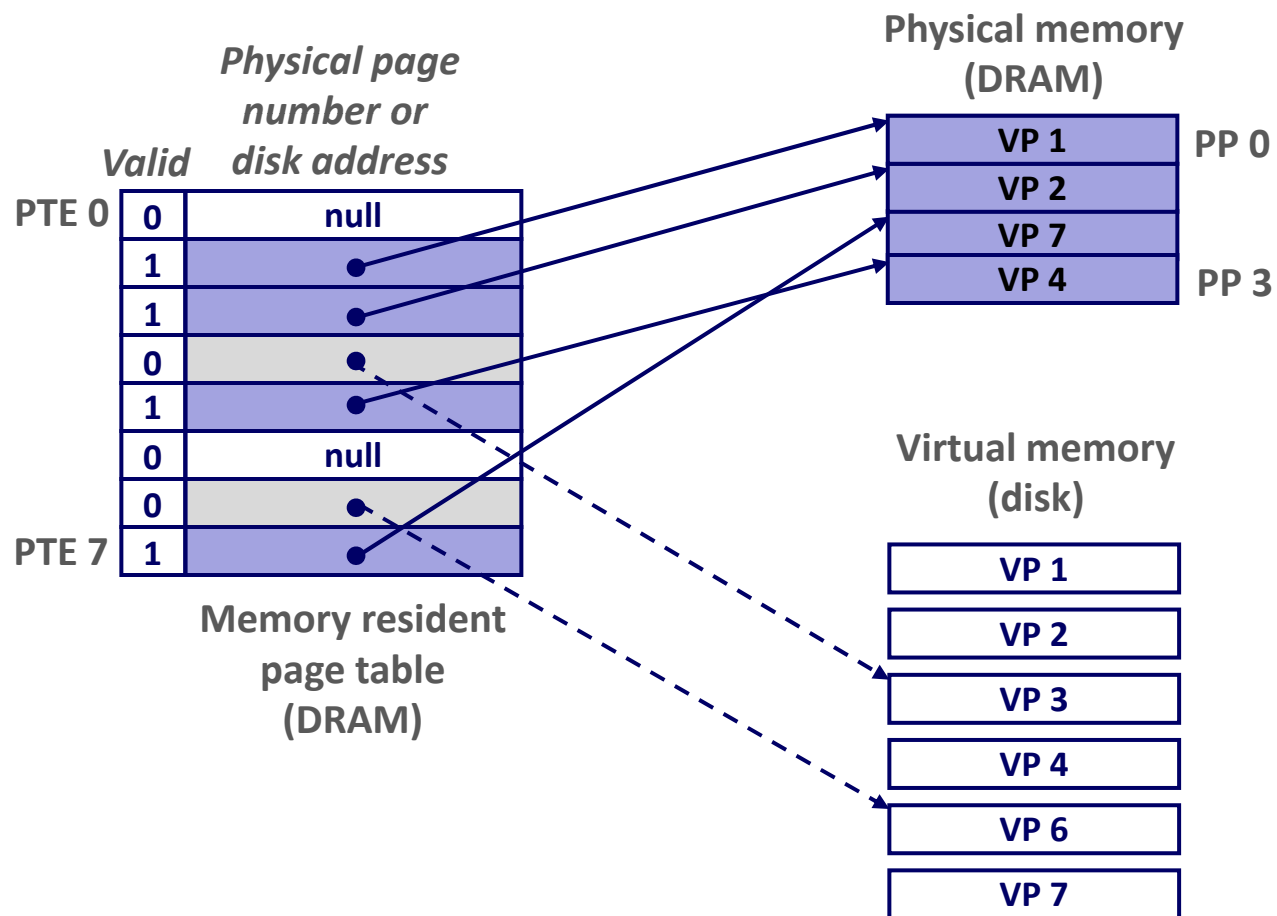
- Used in all modern desktops, laptops, workstations
- One of the great ideas in computer science
- MMU checks the cache

Why Virtual Memory (VM)?

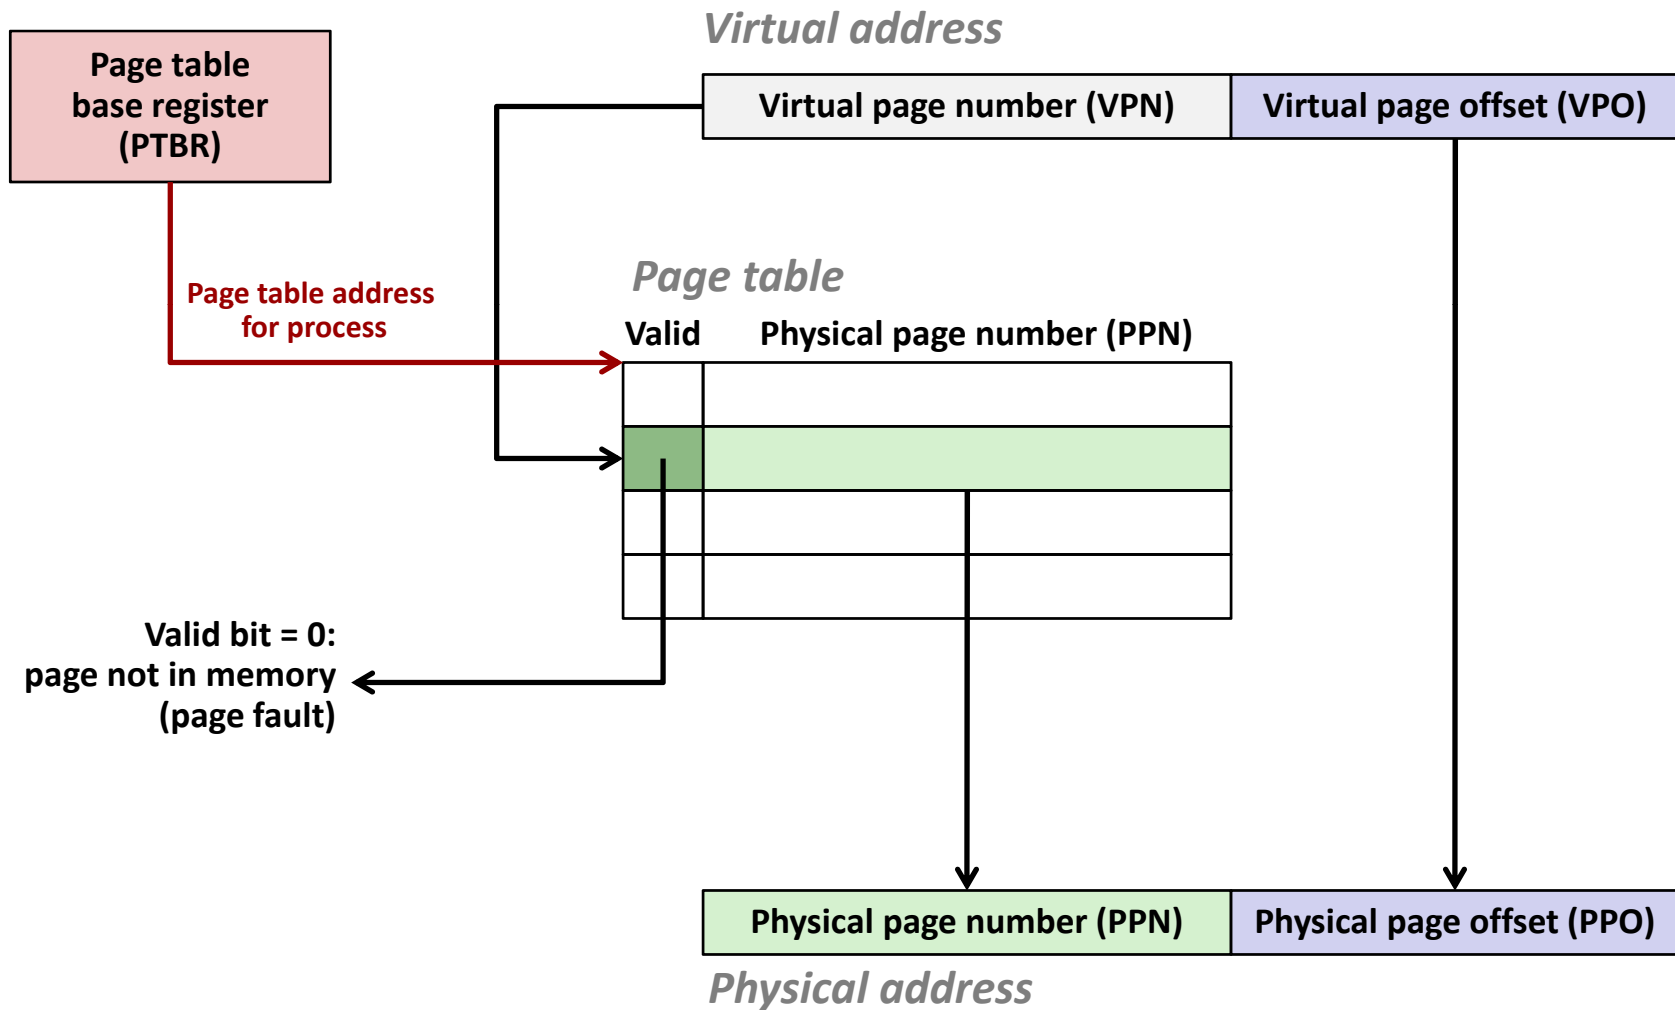
- Efficient use of limited main memory (RAM)
 - Use RAM as a cache for the parts of a virtual address space
 - some non-cached parts stored on disk
 - some (unallocated) non-cached parts stored nowhere
 - Keep only active areas of virtual address space in memory
 - transfer data back and forth as needed
- Simplifies memory management for programmers
 - Each process gets the same full, private linear address space
- Isolates address spaces
 - One process can't interfere with another's memory
 - because they operate in different address spaces
 - User process cannot access privileged information
 - different sections of address spaces have different permissions

Address Translation: Page Tables

- A page table is an array of page table entries (PTEs) that maps virtual pages to physical pages. Here: 8 VPs
 - Per-process kernel data structure in DRAM



Address Translation With a Page Table




Exceptions

Some notes adopted from Bryant and O'Hallaron

Control Flow

- Computers do Only One Thing
 - From startup to shutdown, a CPU simply reads and executes (interprets) a sequence of instructions
 - This sequence is the system's physical *control flow* (or *flow of control*).



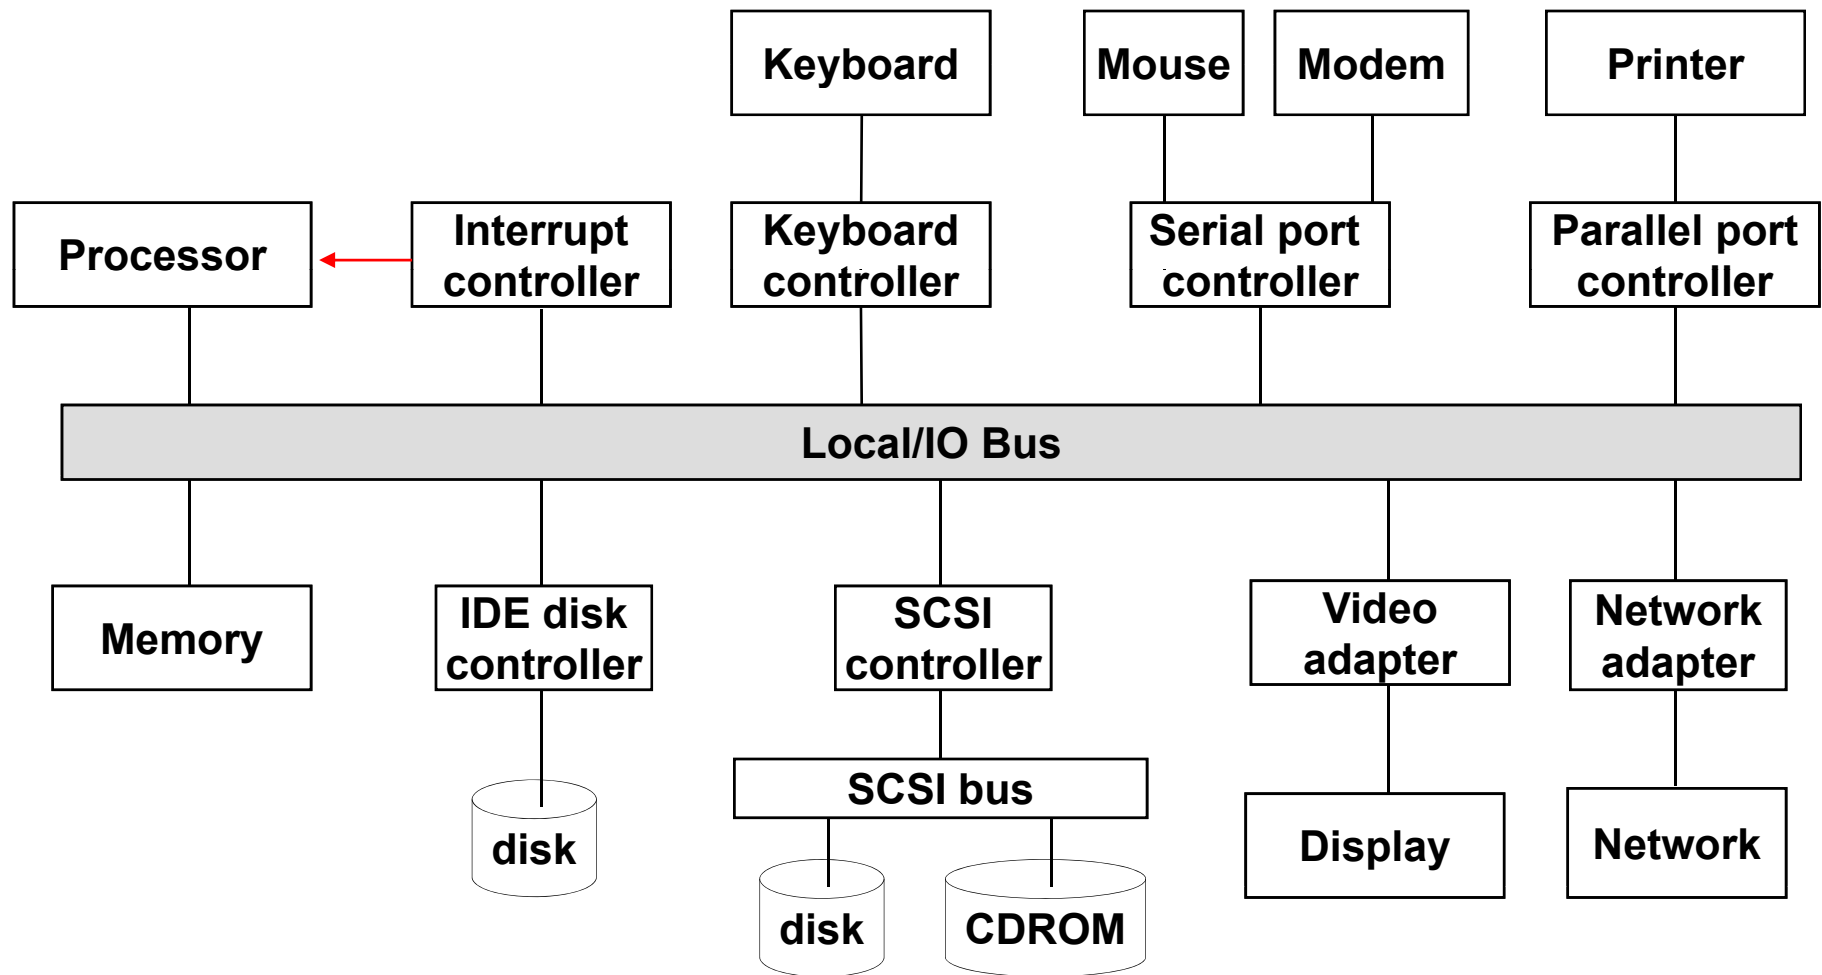
Altering the Control Flow

- Up to Now: two mechanisms for changing control flow:
 - Jumps and branches
 - Call and return using the stack discipline.
 - Both react to changes in program state.
- Insufficient for a useful system
 - Difficult for the CPU to react to changes in system state.
 - data arrives from a disk or a network adapter.
 - Instruction divides by zero
 - User hits ctrl-c at the keyboard
 - System timer expires
- System needs mechanisms for “exceptional control flow”

Exceptional Control Flow

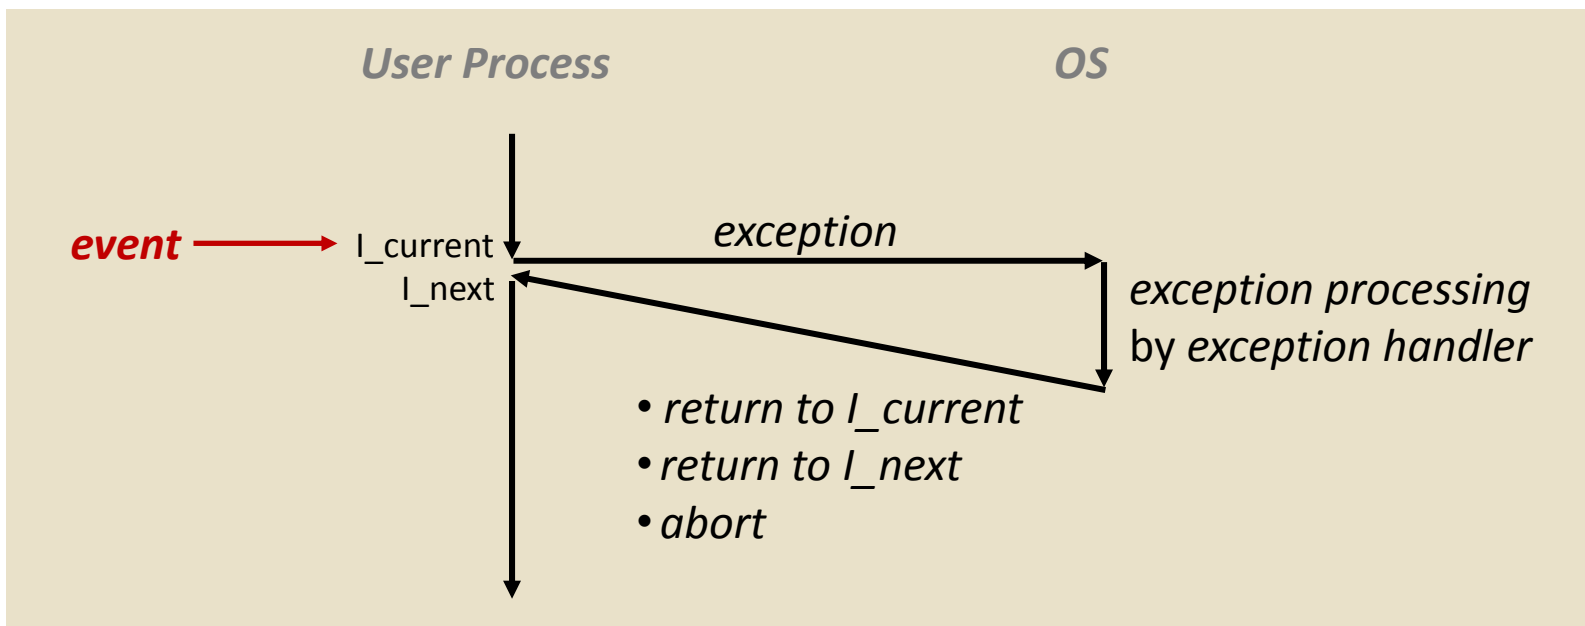
- Mechanisms for exceptional control flow exists at all levels of a computer system.
- Low level Mechanism
 - exceptions
 - change in control flow in response to a system event (i.e., change in system state)
 - Combination of hardware and OS software
- Higher Level Mechanisms
 - Process context switch
 - Signals
 - Nonlocal jumps (setjmp/longjmp)
 - Implemented by either:
 - OS software (context switch and signals).
 - C language runtime library: nonlocal jumps.

System context for exceptions



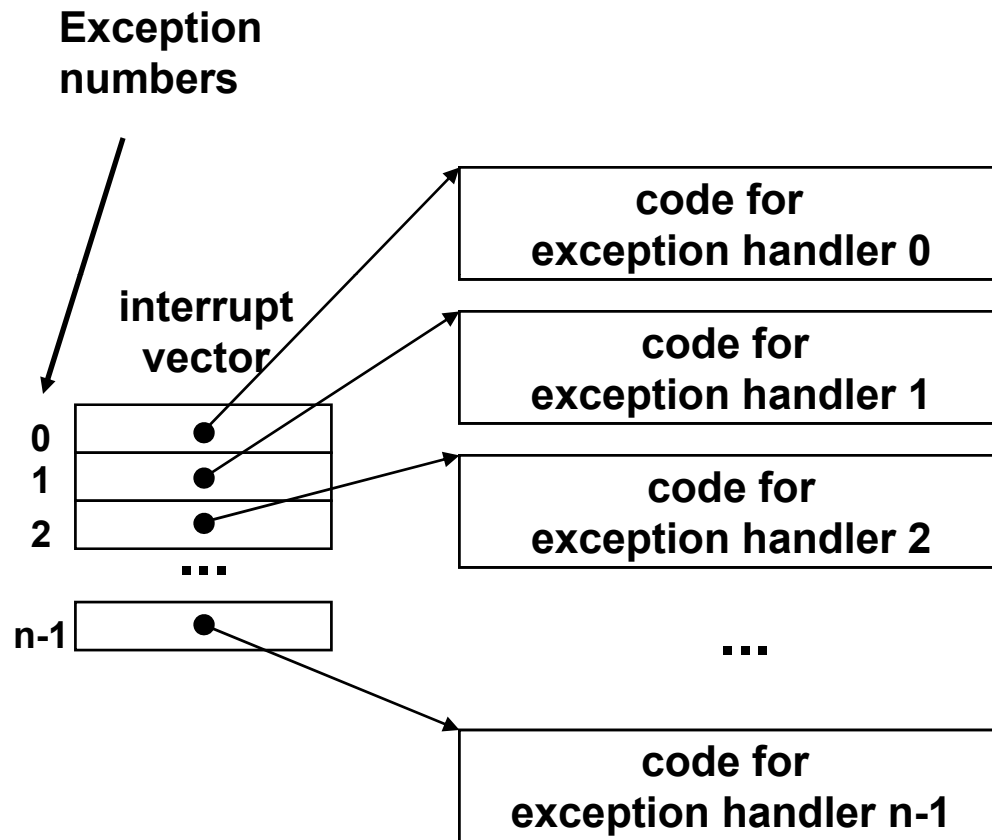
Exceptions

- An exception is a transfer of control to the OS in response to some event (i.e., change in processor state)



- Examples:
div by 0, arithmetic overflow, page fault, I/O request completes, Ctrl-C

Interrupt Vectors

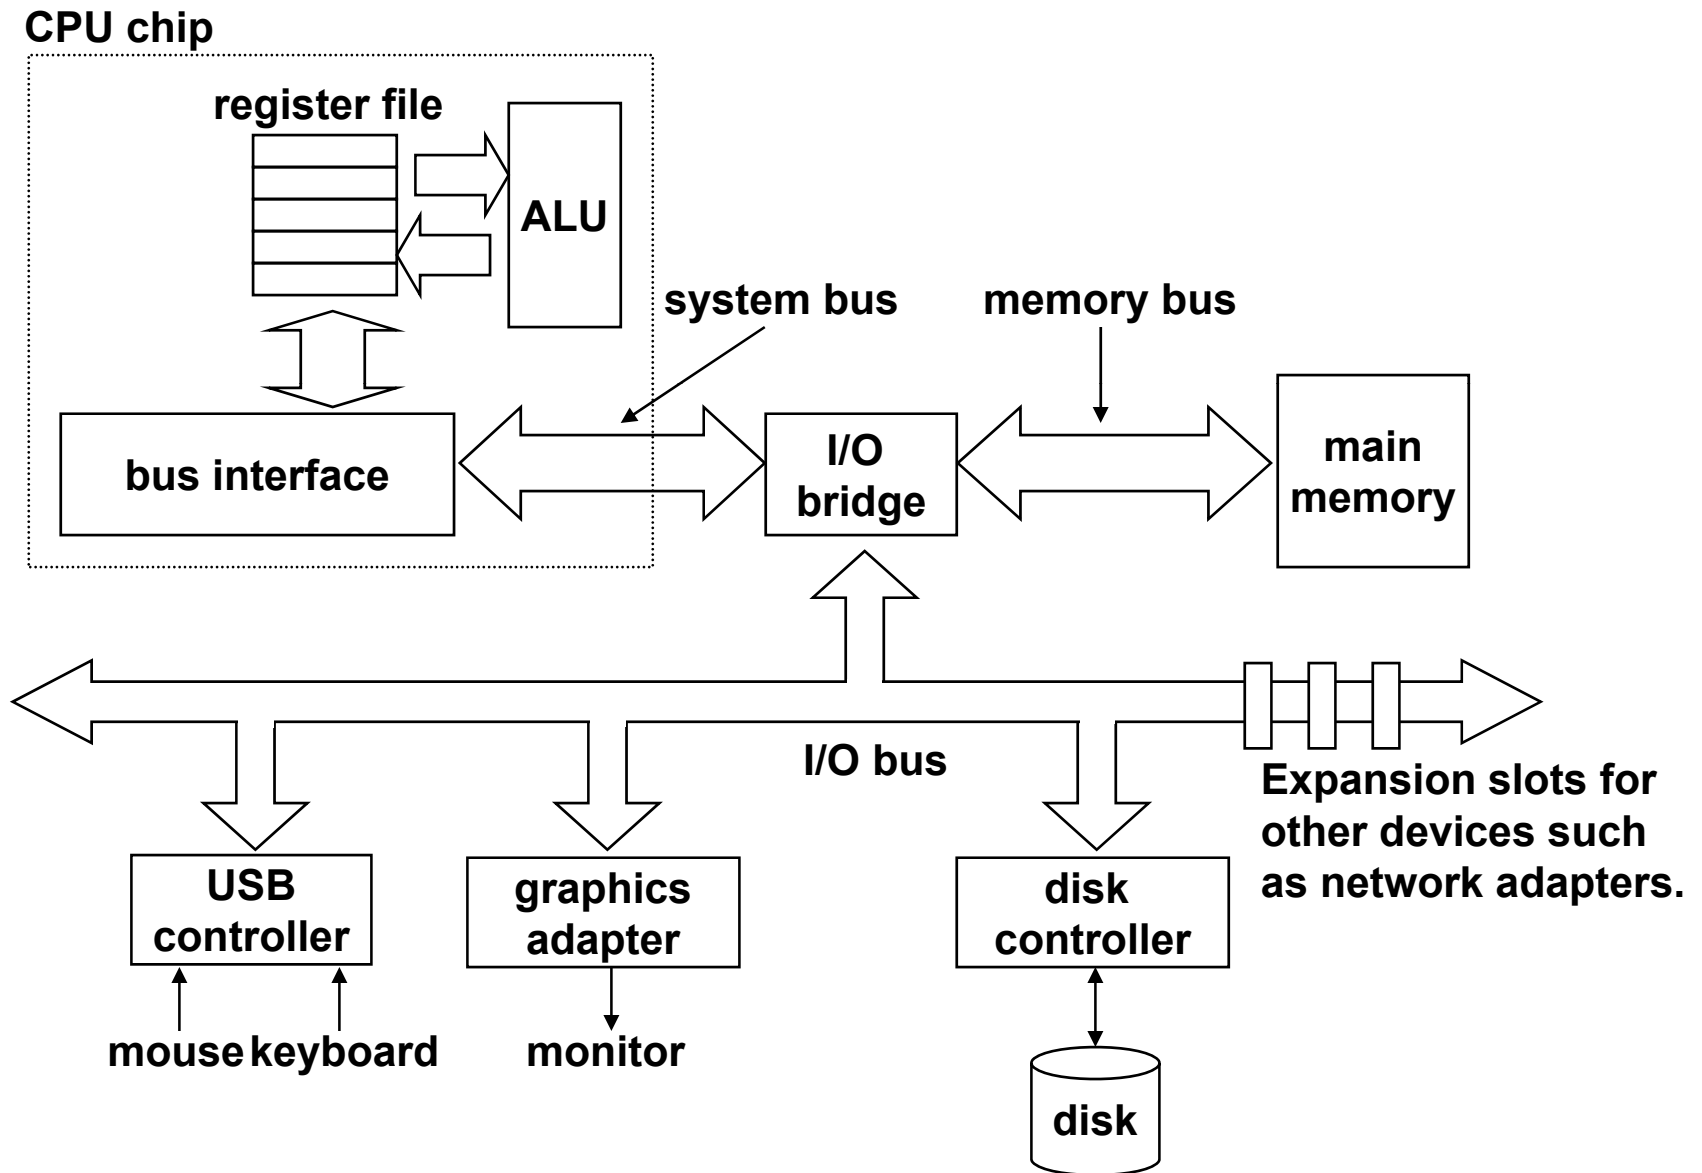


- Each type of event has a unique exception number k
- Index into jump table (a.k.a., interrupt vector)
- Jump table entry k points to a function (exception handler).
- Handler k is called each time exception k occurs.

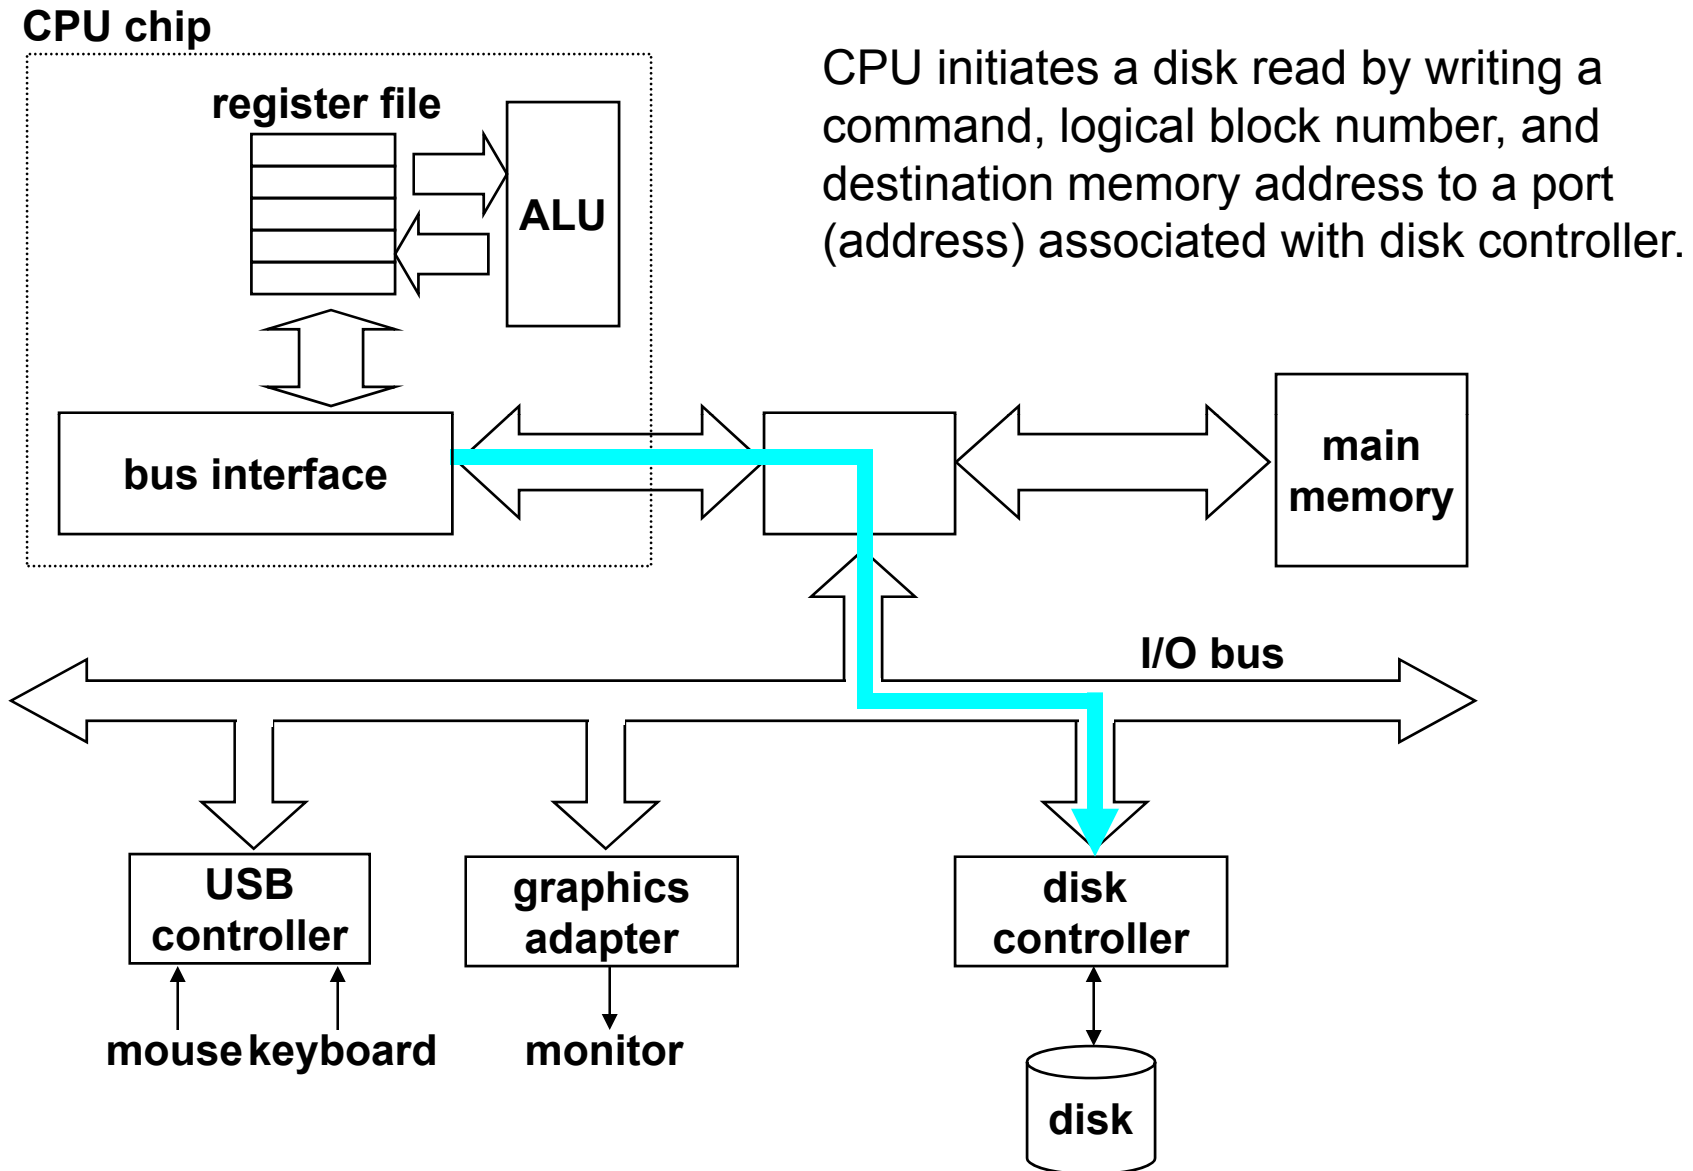
Asynchronous Exceptions (Interrupts)

- Caused by events external to the processor
 - Indicated by setting the processor's interrupt pin
 - handler returns to “next” instruction.
- Examples:
 - I/O interrupts
 - hitting ctrl-c at the keyboard
 - arrival of a packet from a network
 - arrival of a data sector from a disk
 - Hard reset interrupt
 - hitting the reset button
 - Soft reset interrupt
 - hitting ctrl-alt-delete on a PC

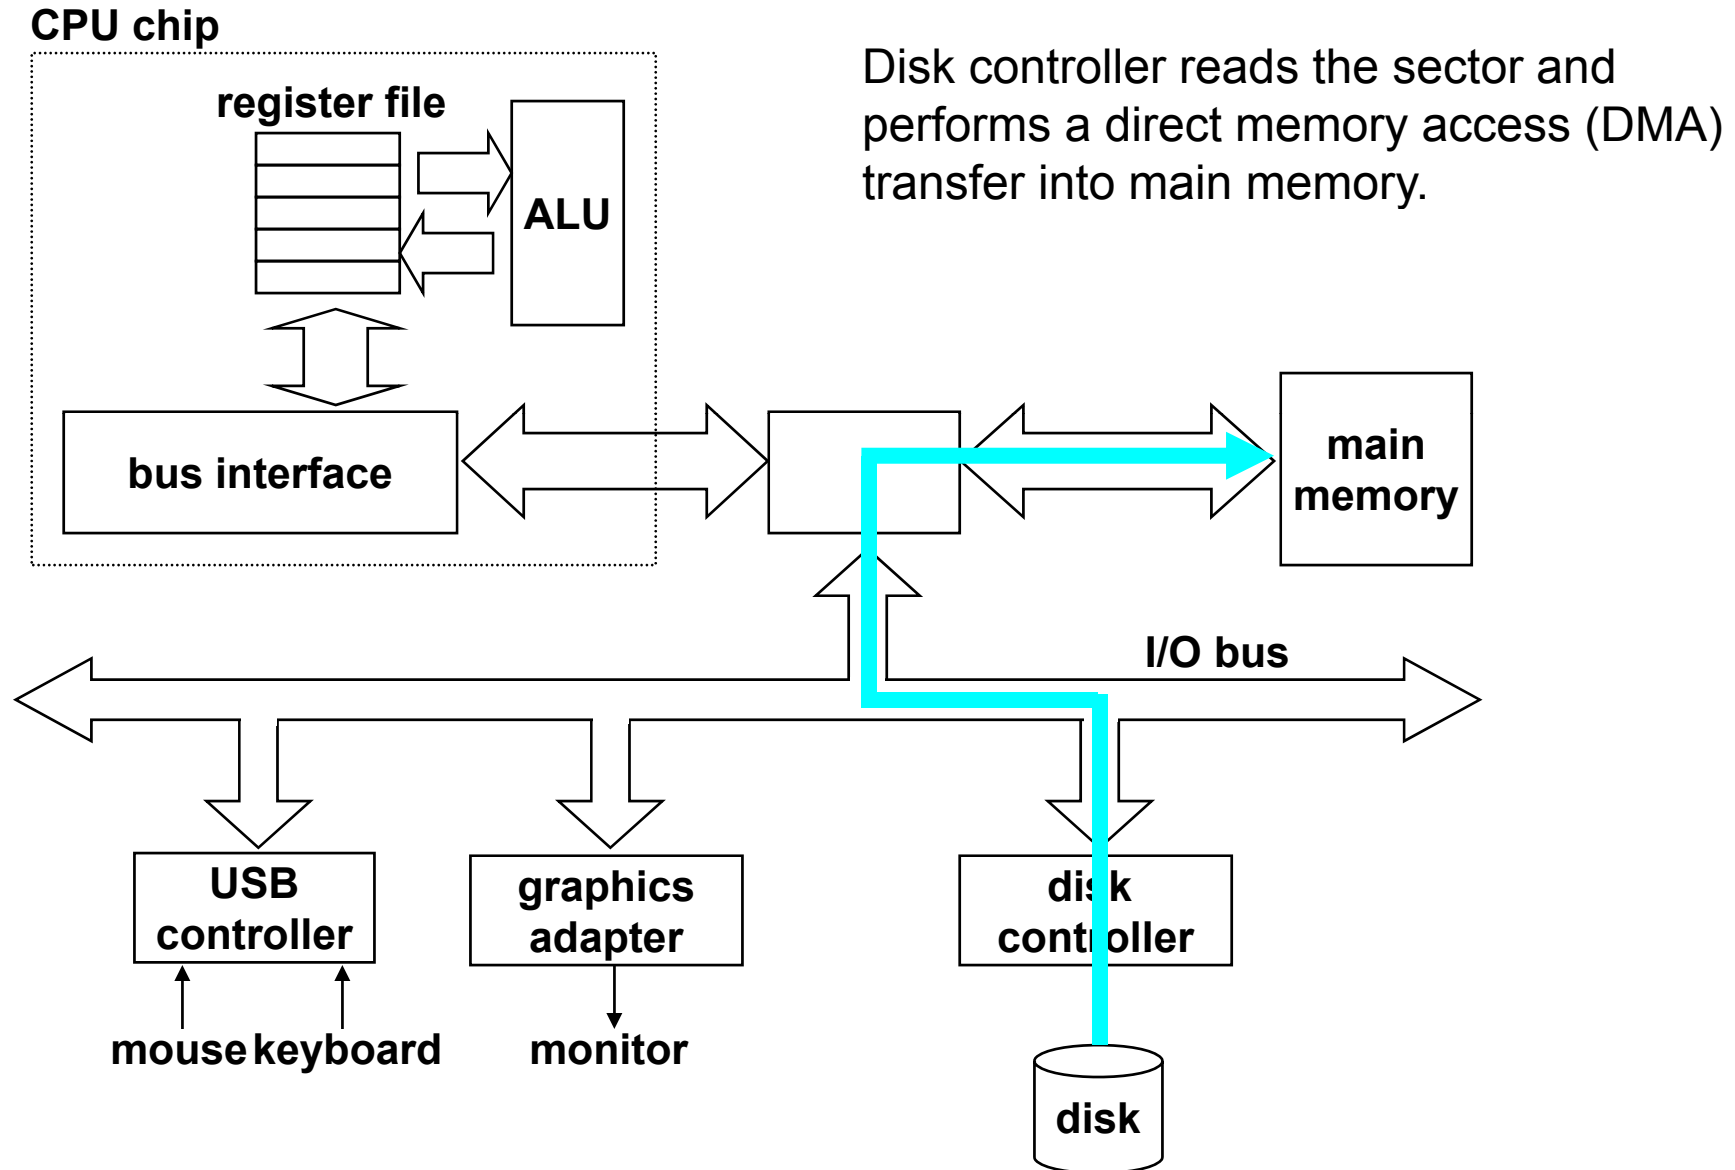
A Typical Hardware System



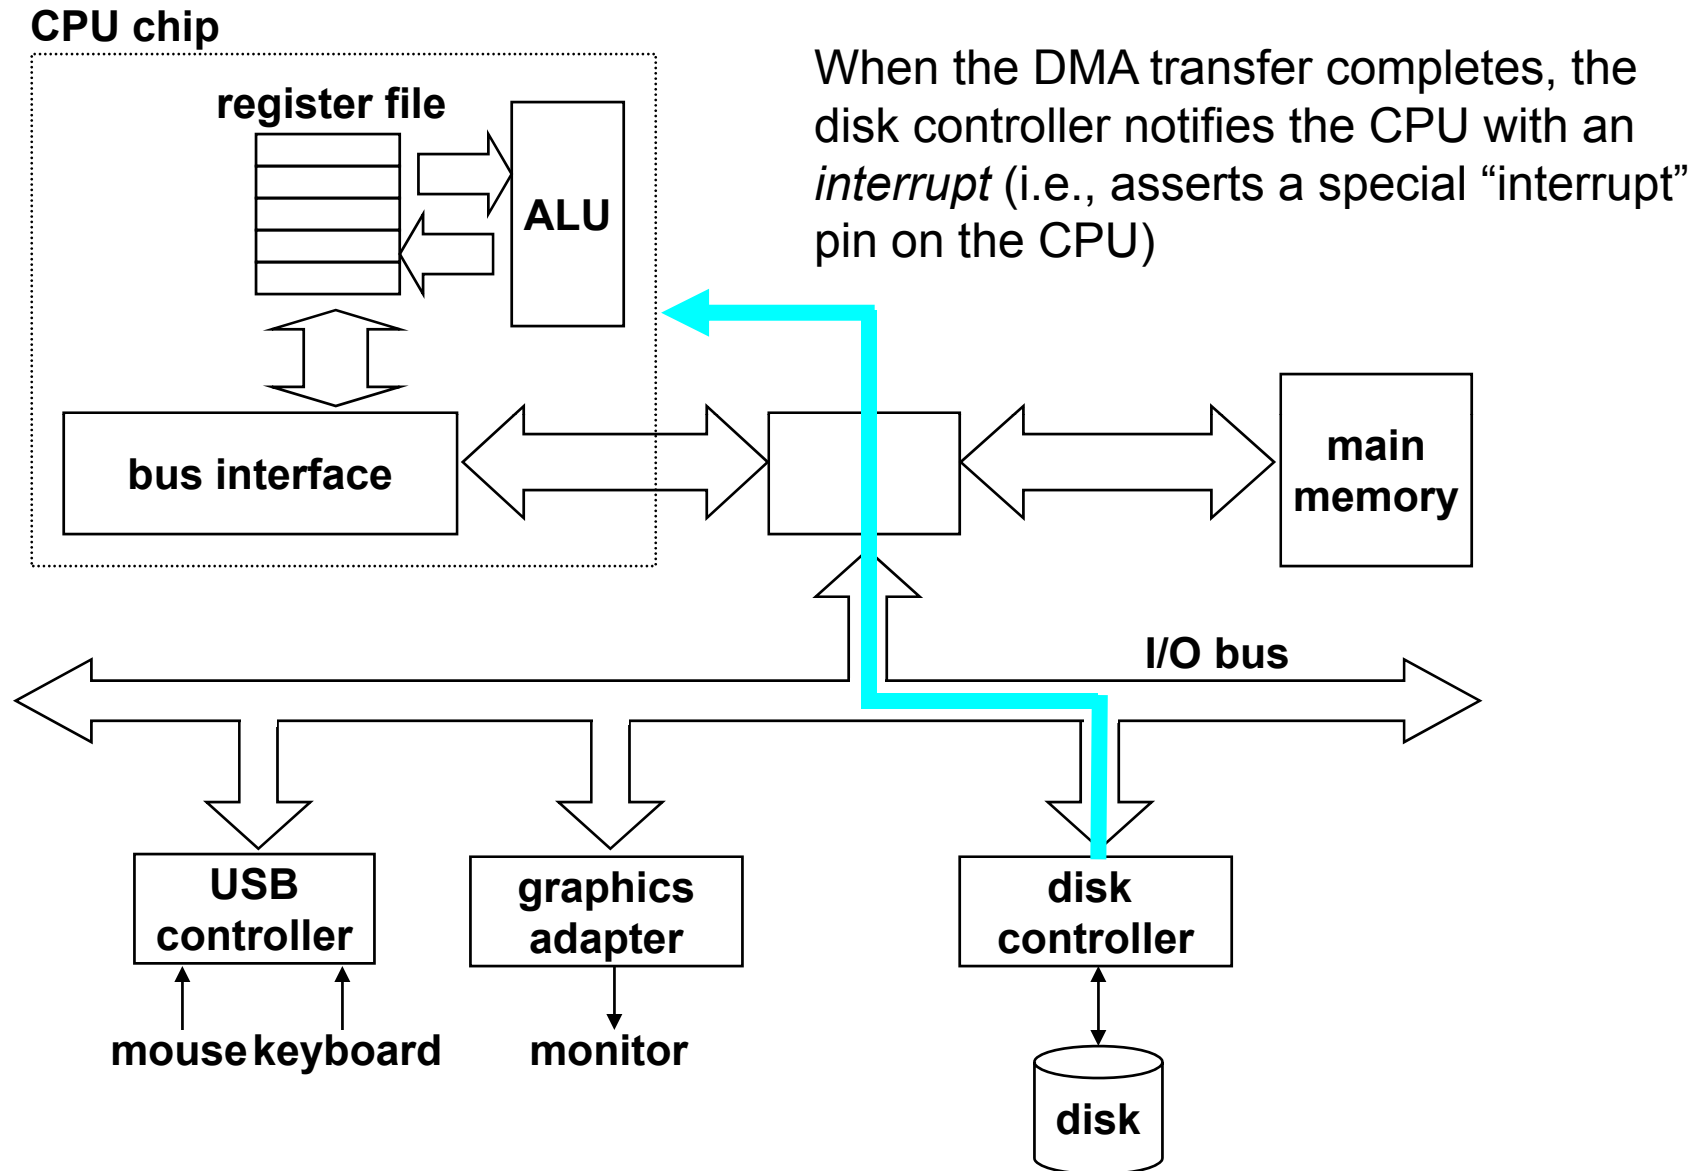
Reading a Disk Sector: Step 1



Reading a Disk Sector: Step 2



Reading a Disk Sector: Step 3



Synchronous Exceptions

- Caused by events that occur as a result of executing an instruction:
 - Traps
 - Intentional
 - Examples: system calls, breakpoint traps, special instructions
 - Returns control to “next” instruction
 - Faults
 - Unintentional but possibly recoverable
 - Examples: page faults (recoverable), protection faults (unrecoverable).
 - Either re-executes faulting (“current”) instruction or aborts.
 - Aborts
 - unintentional and unrecoverable
 - Examples: parity error, machine check.
 - Aborts current program


Trap Example

- Opening a File

- User calls `open(filename, options)`

```
0804d070 <__libc_open>:  
. . .  
804d082:      cd 80                int    $0x80  
804d084:      5b                    pop     %ebx  
. . .
```

- Function `open` executes system call instruction `int`
- OS must find or create file, get it ready for reading or writing
- Returns integer file descriptor

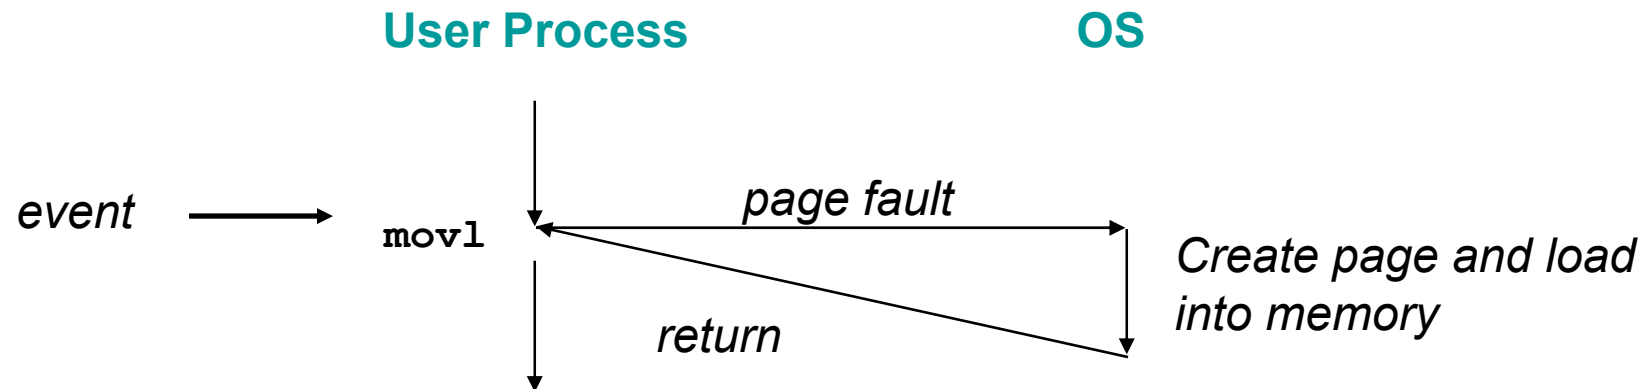


Fault Example #1

- Memory Reference

- User writes to memory location
- That portion (page) of user's memory is currently on disk
- Page handler must load page into physical memory
- Returns to faulting instruction
- Successful on second try

```
int a[1000];  
main ()  
{  
    a[500] = 13;  
}
```

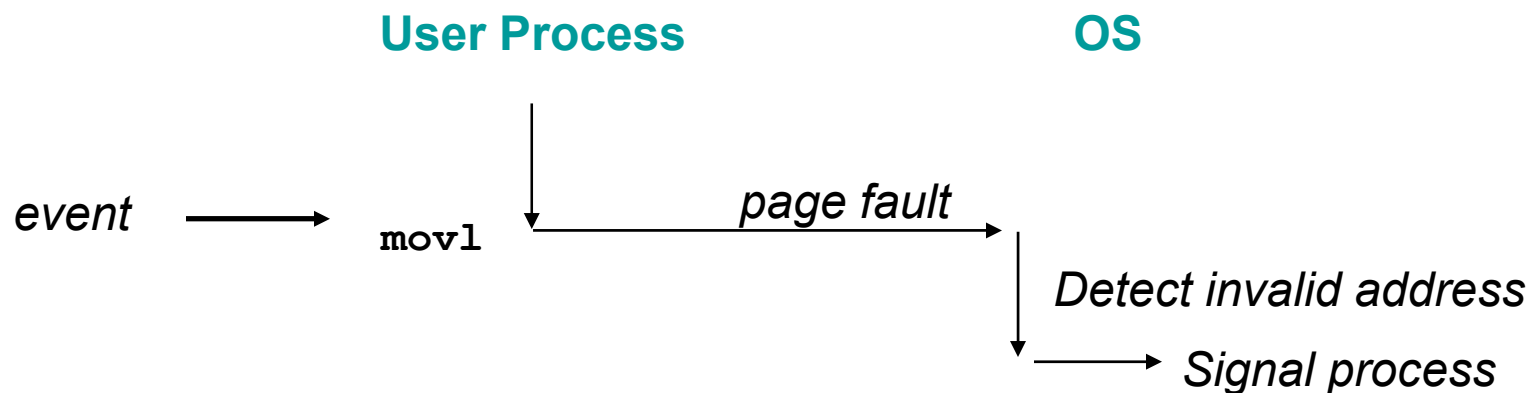


Fault Example #2

- Memory Reference

- User writes to memory location
- Address is not valid
- Page handler detects invalid address
- Sends SIGSEGV signal to user process
- User process exits with “segmentation fault”

```
int a[1000];  
main ()  
{  
    a[5000] = 13;  
}
```



Exception Table IA32 (Excerpt)

<i>Exception Number</i>	<i>Description</i>	<i>Exception Class</i>
0	Divide error	Fault
13	General protection fault	Fault
14	Page fault	Fault
18	Machine check	Abort
32-127	OS-defined	Interrupt or trap
128 (0x80)	System call	Trap
129-255	OS-defined	Interrupt or trap

Check pp. 183:

<http://download.intel.com/design/processor/manuals/253665.pdf>

Linkers

Some notes adopted from Bryant and O'Hallaron

Example C Program

main.c

```
int buf[2] = {1, 2};

int main()
{
    swap();
    return 0;
}
```

swap.c

```
extern int buf[];

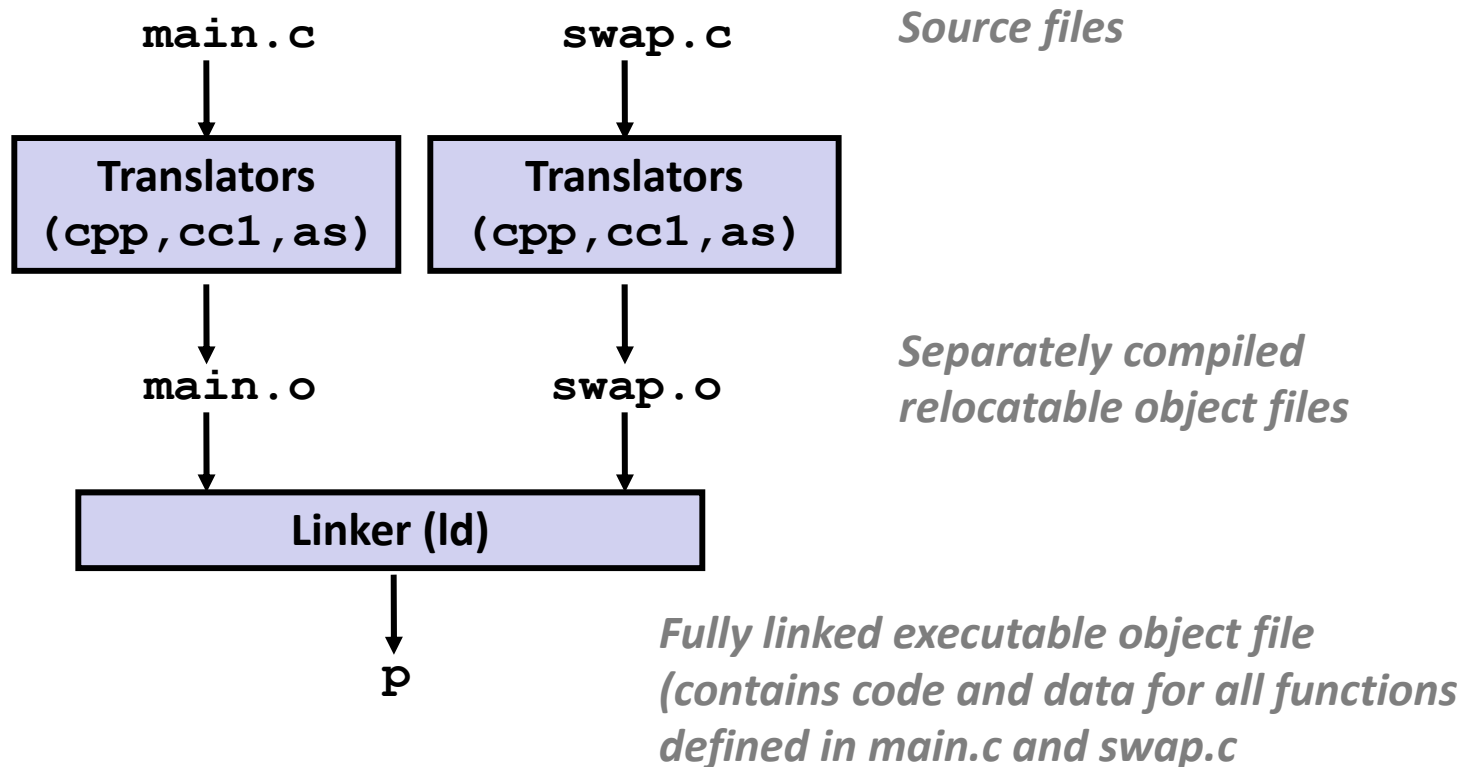
static int *bufp0 = &buf[0];
static int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

Static Linking

- Programs are translated and linked using a compiler driver:
 - `unix> gcc -O2 -g -o p main.c swap.c`
 - `unix> ./p`



Why Linkers? Modularity!

- Program can be written as a collection of smaller source files, rather than one monolithic mass.
- Can build libraries of common functions (more on this later)
 - e.g., Math library, standard C library

Why Linkers? Efficiency!

- Time: Separate Compilation

- Change one source file, compile, and then relink.
- No need to recompile other source files.

- Space: Libraries

- Common functions can be aggregated into a single file...
- Yet executable files and running memory images contain only code for the functions they actually use.

What Do Linkers Do?

- Step 1: Symbol resolution

- Programs define and reference symbols (variables and functions):

- `void swap() {...} /* define symbol swap */`

- `swap(); /* reference symbol swap */`

- `int *xp = &x; /* define xp, reference x */`

- Symbol definitions are stored (by compiler) in symbol table.

- Symbol table is an array of structs

- Each entry includes name, type, size, and location of symbol.

- Linker associates each symbol reference with exactly one symbol definition.

What Do Linkers Do? (cont.)

- Step 2: Relocation

- Merges separate code and data sections into single sections
- Relocates symbols from their relative locations in the .o files to their final absolute memory locations in the executable.
- Updates all references to these symbols to reflect their new positions.

Three Kinds of Object Files (Modules)

- Relocatable object file (.o file)
 - Contains code and data in a form that can be combined with other relocatable object files to form executable object file.
 - Each .o file is produced from exactly one source (.c) file
- Executable object file
 - Contains code and data in a form that can be copied directly into memory and then executed.
- Shared object file (.so file)
 - Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time.
 - Called Dynamic Link Libraries (DLLs) by Windows

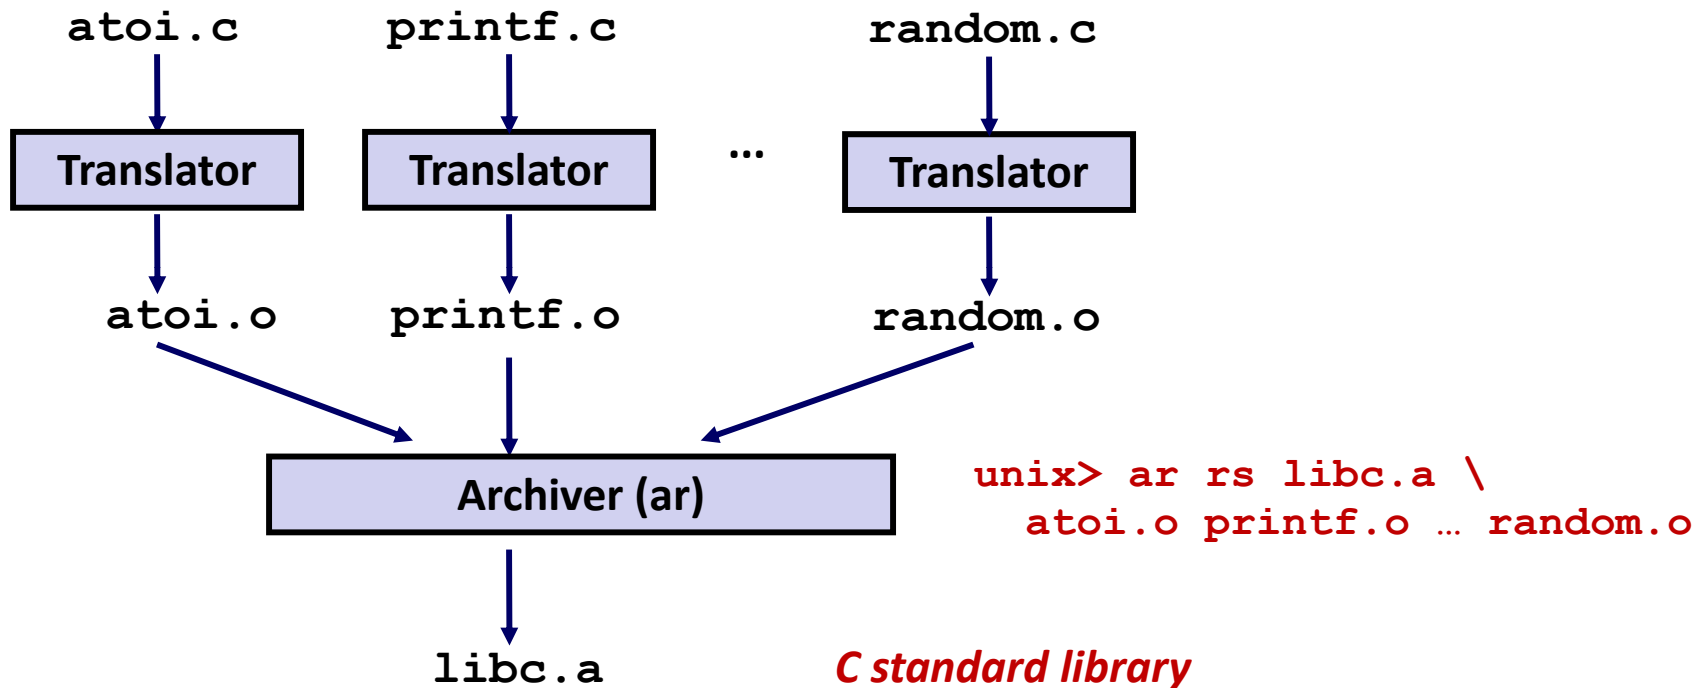
Packaging Commonly Used Functions`

- How to package functions commonly used by programmers?
 - Math, I/O, memory management, string manipulation, etc.
- Awkward, given the linker framework so far:
 - Option 1: Put all functions into a single source file
 - Programmers link big object file into their programs
 - Space and time inefficient
 - Option 2: Put each function in a separate source file
 - Programmers explicitly link appropriate binaries into their programs
 - More efficient, but burdensome on the programmer

Solution: Static Libraries

- Static libraries (.a archive files)
 - Concatenate related relocatable object files into a single file with an index (called an archive).
 - Enhance linker so that it tries to resolve unresolved external references by looking for the symbols in one or more archives.
 - If an archive member file resolves reference, link into executable.

Creating Static Libraries



- Archiver allows incremental updates
- Recompile function that changes and replace .o file in archive.

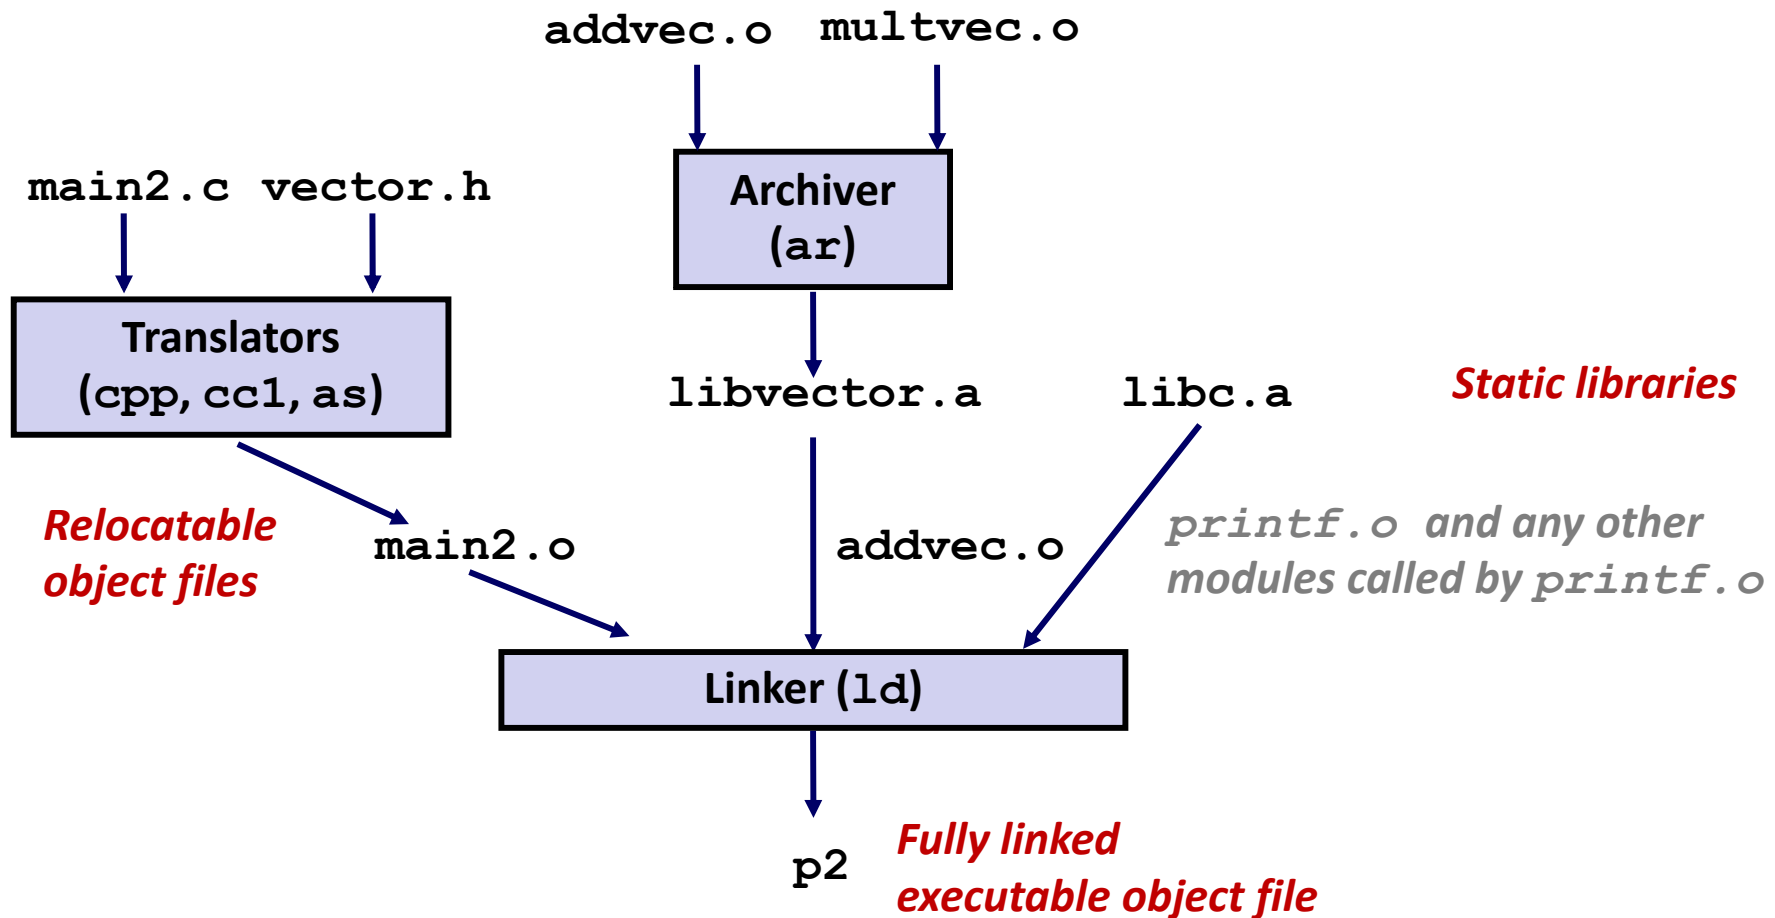
Commonly Used Libraries

- **libc.a (the C standard library)**
 - 8 MB archive of 900 object files.
 - I/O, memory allocation, signal handling, string handling, data and time, random numbers, integer math
- **libm.a (the C math library)**
 - 1 MB archive of 226 object files.
 - floating point math (sin, cos, tan, log, exp, sqrt, ...)


```
% ar -t /usr/lib/libc.a | sort
...
fork.o
...
fprintf.o
fpu_control.o
fputc.o
freopen.o
fscanf.o
fseek.o
fstab.o
...
```

```
% ar -t /usr/lib/libm.a | sort
...
e_acos.o
e_acosf.o
e_acosh.o
e_acoshf.o
e_acoshl.o
e_acosl.o
e_asin.o
e_asinf.o
e_asinl.o
...
```

Linking with Static Libraries



Using Static Libraries

- Linker's algorithm for resolving external references:
 - Scan .o files and .a files in the command line order.
 - During the scan, keep a list of the current unresolved references.
 - As each new .o or .a file, obj, is encountered, try to resolve each unresolved reference in the list against the symbols defined in obj.
 - If any entries in the unresolved list at end of scan, then error.
- Problem: 
- Command line order matters!
- Moral: put libraries at the end of the command line.

```
unix> gcc -L. libtest.o -lmine
unix> gcc -L. -lmine libtest.o
libtest.o: In function `main':
libtest.o(.text+0x4): undefined reference to `libfun'
```

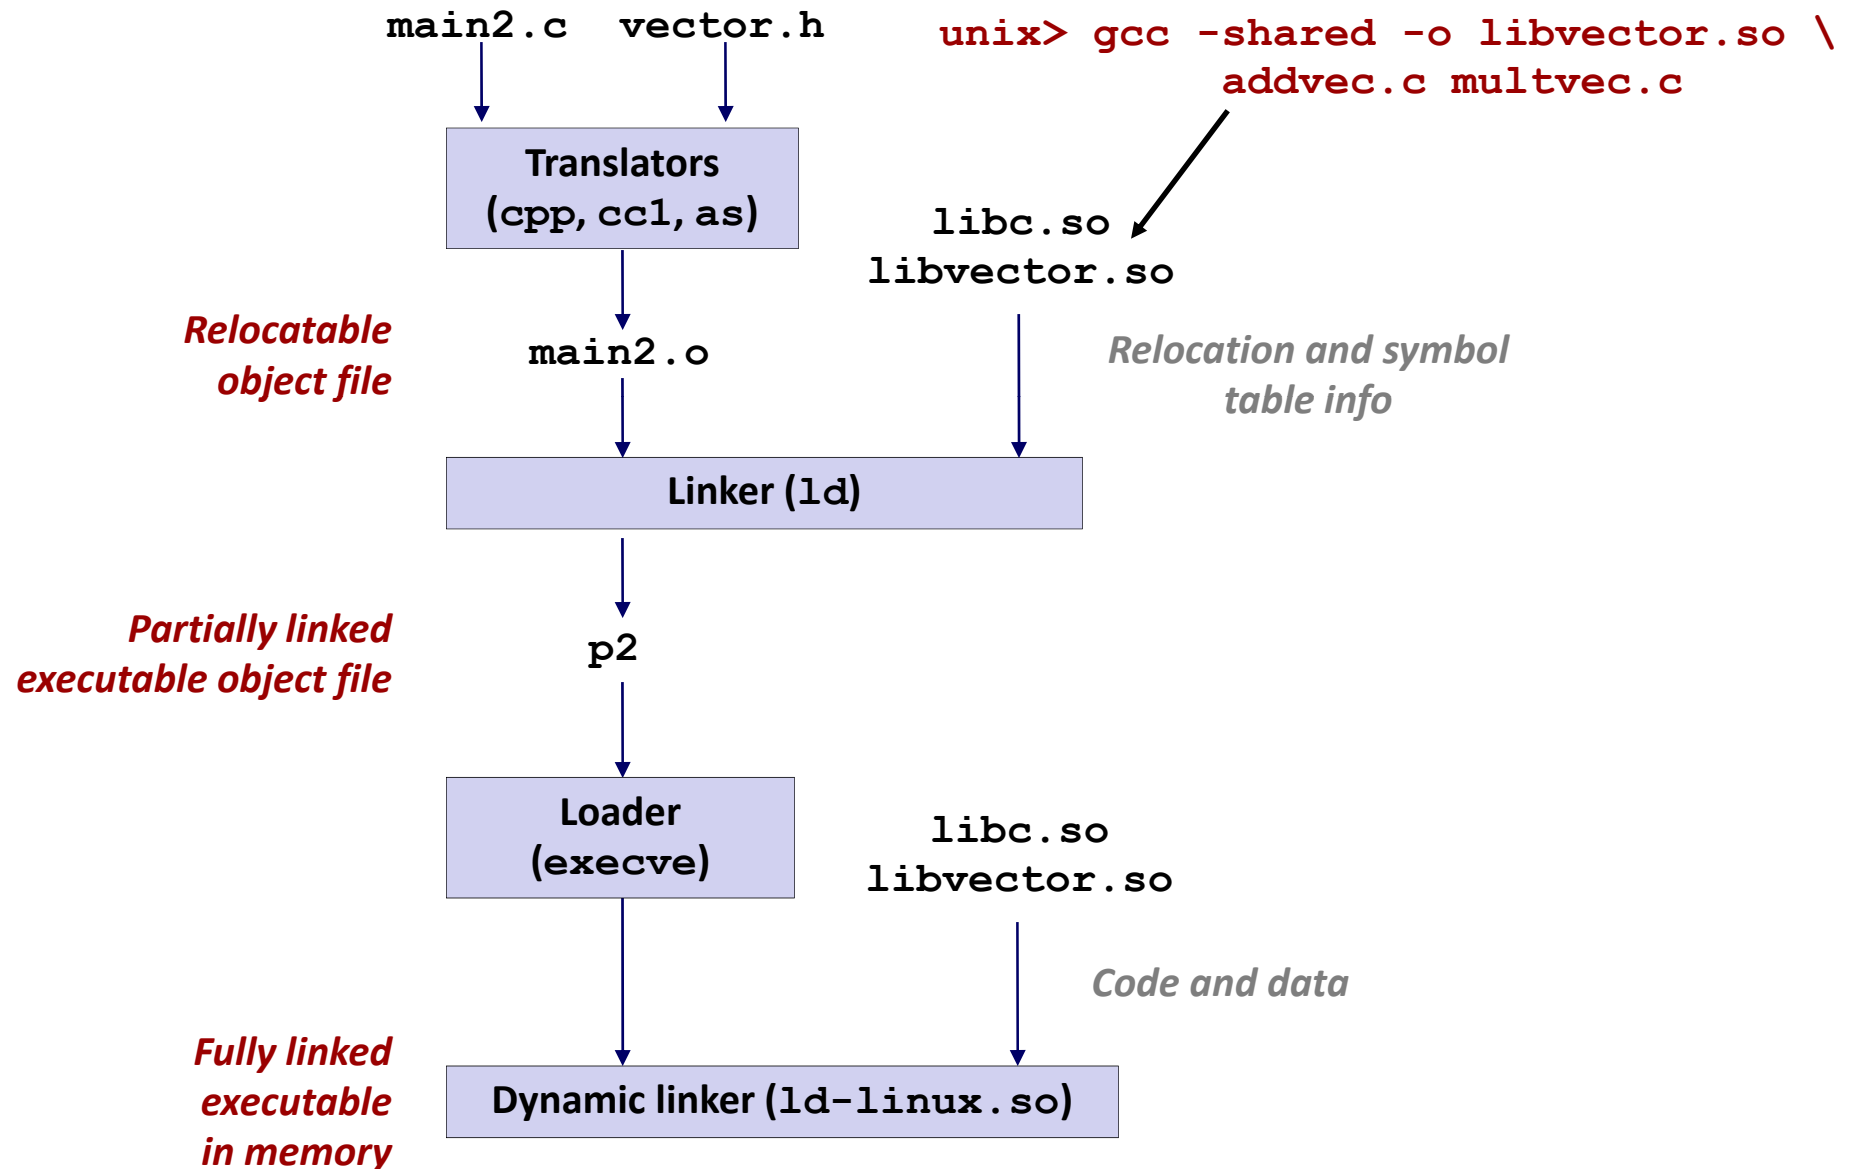
Shared Libraries

- Static libraries have the following disadvantages:
 - Duplication in the stored executables (every function need std libc)
 - Duplication in the running executables
 - Minor bug fixes of system libraries require each application to explicitly relink
- Modern Solution: Shared Libraries
 - Object files that contain code and data that are loaded and linked into an application dynamically, at either load-time or run-time
 - Also called: dynamic link libraries, DLLs, .so files

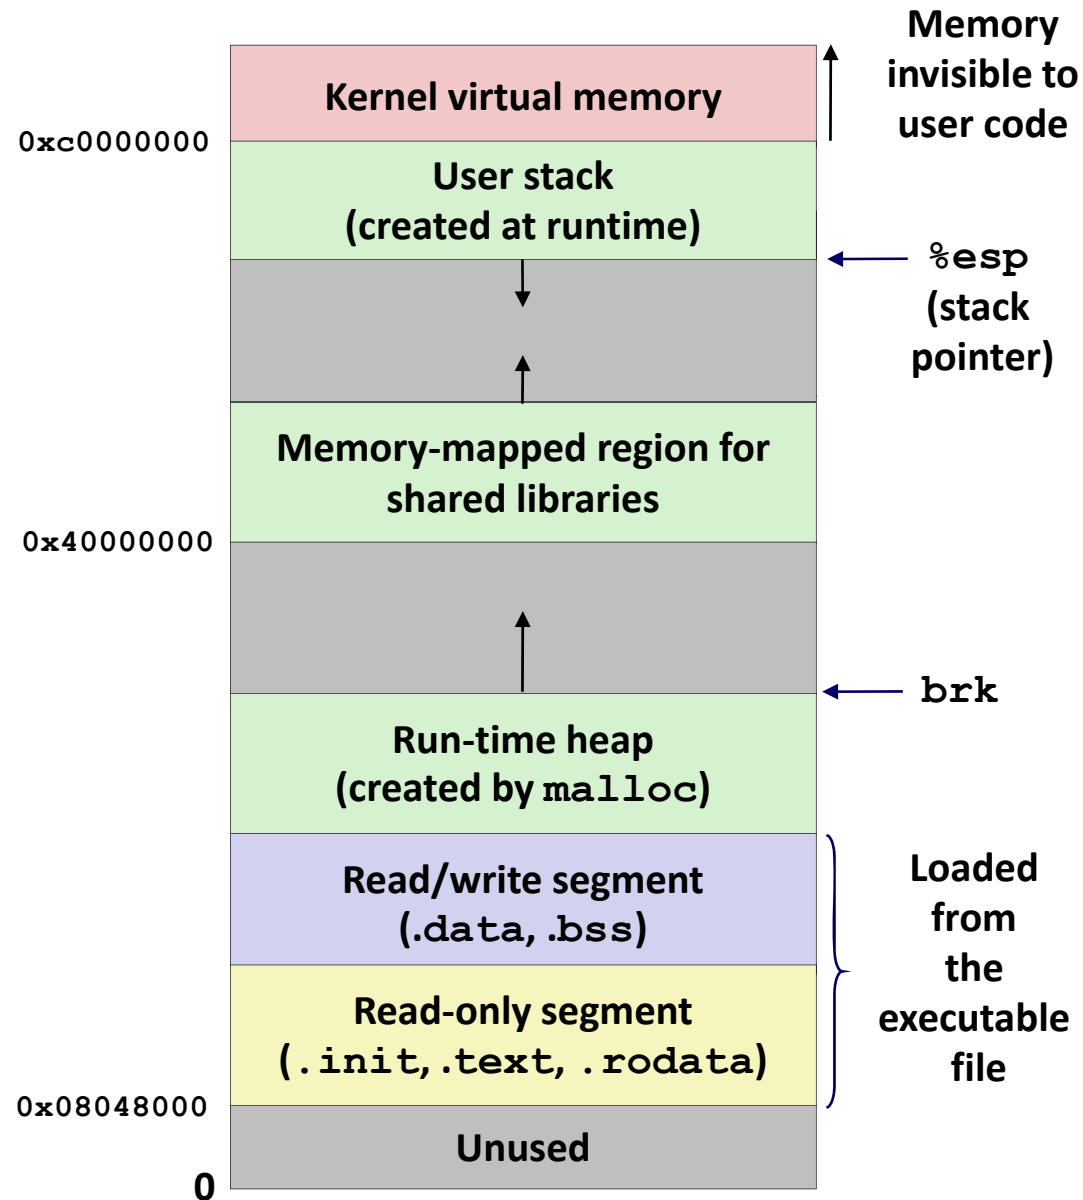
Shared Libraries (cont.)

- Dynamic linking can occur when executable is first loaded and run (load-time linking).
 - Common case for Linux, handled automatically by the dynamic linker (ld-linux.so).
 - Standard C library (libc.so) usually dynamically linked.
- Dynamic linking can also occur after program has begun (run-time linking).
 - In Unix, this is done by calls to the dlopen() interface.
 - High-performance web servers.
 - Runtime library interpositioning
- Shared library routines can be shared by multiple processes.
 - More on this when we learn about virtual memory

Dynamic Linking at Load-time



Refined View of Memory



Case Study: Library Interpositioning

- Library interpositioning is a powerful linking technique that allows programmers to intercept calls to arbitrary functions
- Interpositioning can occur at:
 - compile time
 - When the source code is compiled
 - link time
 - When the relocatable object files are linked to form an executable object file
 - load/run time
 - When an executable object file is loaded into memory, dynamically linked, and then executed.

Some Interpositioning Applications

- Security

- Confinement (sandboxing)
 - Interpose calls to libc functions.
- Behind the scenes encryption
 - Automatically encrypt otherwise unencrypted network connections.

- Monitoring and Profiling

- Count number of calls to functions
- Characterize call sites and arguments to functions
- Malloc tracing
 - Detecting memory leaks
 - Generating malloc traces