

NVIDIA®

CUDA Basics

CUDA

A Parallel Computing Architecture for NVIDIA GPUs



Development Platform of Choice

- Over 60,000 GPU Computing Developers (1/09)
- Windows, Linux and MacOS Platforms supported
- Mature Development tools


GPU Computing Applications

C
with CUDA extensions

OpenCL™

**DirectX
Compute**

FORTRAN



NVIDIA GPU

with the CUDA Parallel Computing Architecture



Outline of CUDA Basics

- **Basics Memory Management**
- **Basic Kernels and Execution on GPU**
- **Coordinating CPU and GPU Execution**
- **Development Resources**
- **See the Programming Guide for the full API**



NVIDIA®

Basic Memory Management

Memory Spaces

- **CPU and GPU have separate memory spaces**
 - Data is moved across PCIe bus
 - Use functions to allocate/set/copy memory on GPU
 - Very similar to corresponding C functions
- **Pointers are just addresses**
 - Can't tell from the pointer value whether the address is on CPU or GPU
 - Must exercise care when dereferencing:
 - Dereferencing CPU pointer on GPU will likely crash
 - Same for vice versa

GPU Memory Allocation / Release



- **Host (CPU) manages device (GPU) memory:**
 - `cudaMalloc (void ** pointer, size_t nbytes)`
 - `cudaMemset (void * pointer, int value, size_t count)`
 - `cudaFree (void* pointer)`

```
int n = 1024;  
int nbytes = 1024*sizeof(int);  
int * d_a = 0;  
cudaMalloc( (void**)&d_a, nbytes );  
cudaMemset( d_a, 0, nbytes);  
cudaFree(d_a);
```

Data Copies

- **cudaMemcpy(void *dst, void *src, size_t nbytes, enum cudaMemcpyKind direction);**
 - returns after the copy is complete
 - blocks CPU thread until all bytes have been copied
 - doesn't start copying until previous CUDA calls complete
- **enum cudaMemcpyKind**
 - cudaMemcpyHostToDevice
 - cudaMemcpyDeviceToHost
 - cudaMemcpyDeviceToDevice
- **Non-blocking memcpyes are provided**



Code Walkthrough 1

- **Allocate CPU memory for n integers**
- **Allocate GPU memory for n integers**
- **Initialize GPU memory to 0s**
- **Copy from GPU to CPU**
- **Print the values**

Code Walkthrough 1



```
#include <stdio.h>

int main()
{
    int dimx = 16;
    int num_bytes = dimx*sizeof(int);

    int *d_a=0, *h_a=0; // device and host pointers
```



Code Walkthrough 1

```
#include <stdio.h>

int main()
{
    int dimx = 16;
    int num_bytes = dimx*sizeof(int);

    int *d_a=0, *h_a=0; // device and host pointers

    h_a = (int*)malloc(num_bytes);
    cudaMalloc( (void**)&d_a, num_bytes );

    if( 0==h_a || 0==d_a )
    {
        printf("couldn't allocate memory\n");
        return 1;
    }
}
```

Code Walkthrough 1



```
#include <stdio.h>

int main()
{
    int dimx = 16;
    int num_bytes = dimx*sizeof(int);

    int *d_a=0, *h_a=0; // device and host pointers

    h_a = (int*)malloc(num_bytes);
    cudaMalloc( (void**)&d_a, num_bytes );

    if( 0==h_a || 0==d_a )
    {
        printf("couldn't allocate memory\n");
        return 1;
    }

    cudaMemset( d_a, 0, num_bytes );
    cudaMemcpy( h_a, d_a, num_bytes, cudaMemcpyDeviceToHost );
}
```

Code Walkthrough 1



```
#include <stdio.h>

int main()
{
    int dimx = 16;
    int num_bytes = dimx*sizeof(int);

    int *d_a=0, *h_a=0; // device and host pointers

    h_a = (int*)malloc(num_bytes);
    cudaMalloc( (void**)&d_a, num_bytes );

    if( 0==h_a || 0==d_a )
    {
        printf("couldn't allocate memory\n");
        return 1;
    }

    cudaMemset( d_a, 0, num_bytes );
    cudaMemcpy( h_a, d_a, num_bytes, cudaMemcpyDeviceToHost );

    for(int i=0; i<dimx; i++)
        printf("%d ", h_a[i] );
    printf("\n");

    free( h_a );
    cudaFree( d_a );

    return 0;
}
```



NVIDIA®

Basic Kernels and Execution on GPU

CUDA Programming Model



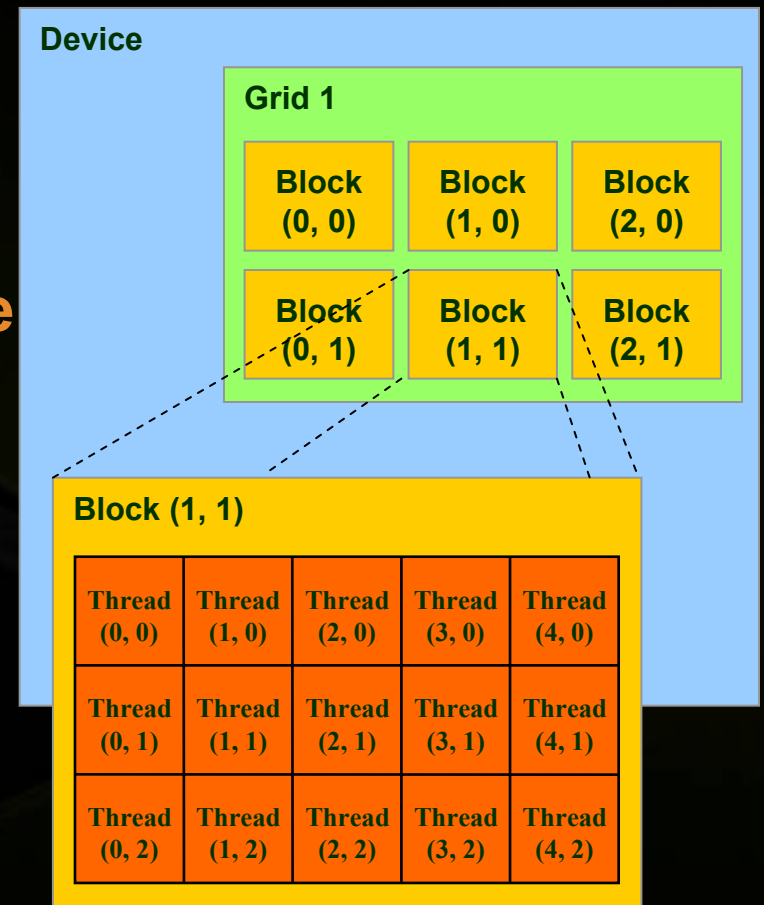
- **Parallel code (kernel) is launched and executed on a device by many threads**
- **Threads are grouped into thread blocks**
- **Parallel code is written for a thread**
 - Each thread is free to execute a unique code path
 - Built-in thread and block ID variables

Thread Hierarchy

- **Threads launched for a parallel section are partitioned into thread blocks**
 - Grid = all blocks for a given launch
- **Thread block is a group of threads that can:**
 - Synchronize their execution
 - Communicate via shared memory

IDs and Dimensions

- **Threads:**
 - 3D IDs, unique within a block
- **Blocks:**
 - 2D IDs, unique within a grid
- **Dimensions set at launch time**
 - Can be unique for each grid
- **Built-in variables:**
 - threadIdx, blockIdx
 - blockDim, gridDim



Code executed on GPU

- **C function with some restrictions:**

- Can only access GPU memory
- No variable number of arguments
- No static variables
- No recursion

- **Must be declared with a qualifier:**

- **__global__** : launched by CPU,
cannot be called from GPU must return void
- **__device__** : called from other GPU functions,
cannot be launched by the CPU
- **__host__** : can be executed by CPU
- **__host__** and **__device__** qualifiers can be combined
 - sample use: overloading operators



Code Walkthrough 2

- **Build on Walkthrough 1**
- **Write a kernel to initialize integers**
- **Copy the result back to CPU**
- **Print the values**

Kernel Code (executed on GPU)



```
__global__ void kernel( int *a )  
{  
    int idx = blockIdx.x*blockDim.x + threadIdx.x;  
    a[idx] = 7;  
}
```



Launching kernels on GPU

● Launch parameters:

- grid dimensions (up to 2D), **dim3** type
- thread-block dimensions (up to 3D), **dim3** type
- shared memory: number of bytes per block
 - for extern smem variables declared without size
 - Optional, 0 by default
- stream ID
 - Optional, 0 by default

```
dim3 grid(16, 16);  
dim3 block(16,16);  
kernel<<<grid, block, 0, 0>>>(...);  
kernel<<<32, 512>>>(...);
```



```
#include <stdio.h>

__global__ void kernel( int *a )
{
    int idx = blockIdx.x*blockDim.x + threadIdx.x;
    a[idx] = 7;
}

int main()
{
    int dimx = 16;
    int num_bytes = dimx*sizeof(int);

    int *d_a=0, *h_a=0; // device and host pointers

    h_a = (int*)malloc(num_bytes);
    cudaMalloc( (void**)&d_a, num_bytes );

    if( 0==h_a || 0==d_a )
    {
        printf("couldn't allocate memory\n");
        return 1;
    }

    cudaMemset( d_a, 0, num_bytes );

    dim3 grid, block;
    block.x = 4;
    grid.x = dimx / block.x;

    kernel<<<grid, block>>>( d_a );

    cudaMemcpy( h_a, d_a, num_bytes, cudaMemcpyDeviceToHost );

    for(int i=0; i<dimx; i++)
        printf("%d ", h_a[i] );
    printf("\n");

    free( h_a );
    cudaFree( d_a );

    return 0;
}
```



Kernel Variations and Output

```
__global__ void kernel( int *a )  
{  
    int idx = blockIdx.x*blockDim.x + threadIdx.x;  
    a[idx] = 7;  
}
```

Output: 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7

```
__global__ void kernel( int *a )  
{  
    int idx = blockIdx.x*blockDim.x + threadIdx.x;  
    a[idx] = blockIdx.x;  
}
```

Output: 0 0 0 0 1 1 1 1 2 2 2 2 3 3 3 3

```
__global__ void kernel( int *a )  
{  
    int idx = blockIdx.x*blockDim.x + threadIdx.x;  
    a[idx] = threadIdx.x;  
}
```

Output: 0 1 2 3 0 1 2 3 0 1 2 3 0 1 2 3



Code Walkthrough 3

- Build on Walkthrough 2
- Write a kernel to increment $n \times m$ integers
- Copy the result back to CPU
- Print the values

Kernel with 2D Indexing



```
__global__ void kernel( int *a, int dimx, int dimy )  
{  
    int ix  = blockIdx.x*blockDim.x + threadIdx.x;  
    int iy  = blockIdx.y*blockDim.y + threadIdx.y;  
    int idx = iy*dimx + ix;  
  
    a[idx] = a[idx]+1;  
}
```




```
__global__ void kernel( int *a, int dimx, int dimy )
{
    int ix  = blockIdx.x*blockDim.x + threadIdx.x;
    int iy  = blockIdx.y*blockDim.y + threadIdx.y;
    int idx = iy*dimx + ix;

    a[idx] = a[idx]+1;
}
```

```
int main()
{
    int dimx = 16;
    int dimy = 16;
    int num_bytes = dimx*dimy*sizeof(int);

    int *d_a=0, *h_a=0; // device and host pointers

    h_a = (int*)malloc(num_bytes);
    cudaMalloc( (void**)&d_a, num_bytes );

    if( 0==h_a || 0==d_a )
    {
        printf("couldn't allocate memory\n");
        return 1;
    }

    cudaMemset( d_a, 0, num_bytes );

    dim3 grid, block;
    block.x = 4;
    block.y = 4;
    grid.x = dimx / block.x;
    grid.y = dimy / block.y;

    kernel<<<grid, block>>>( d_a, dimx, dimy );

    cudaMemcpy( h_a, d_a, num_bytes, cudaMemcpyDeviceToHost );

    for(int row=0; row<dimy; row++)
    {
        for(int col=0; col<dimx; col++)
            printf("%d ", h_a[row*dimx+col] );
        printf("\n");
    }

    free( h_a );
    cudaFree( d_a );

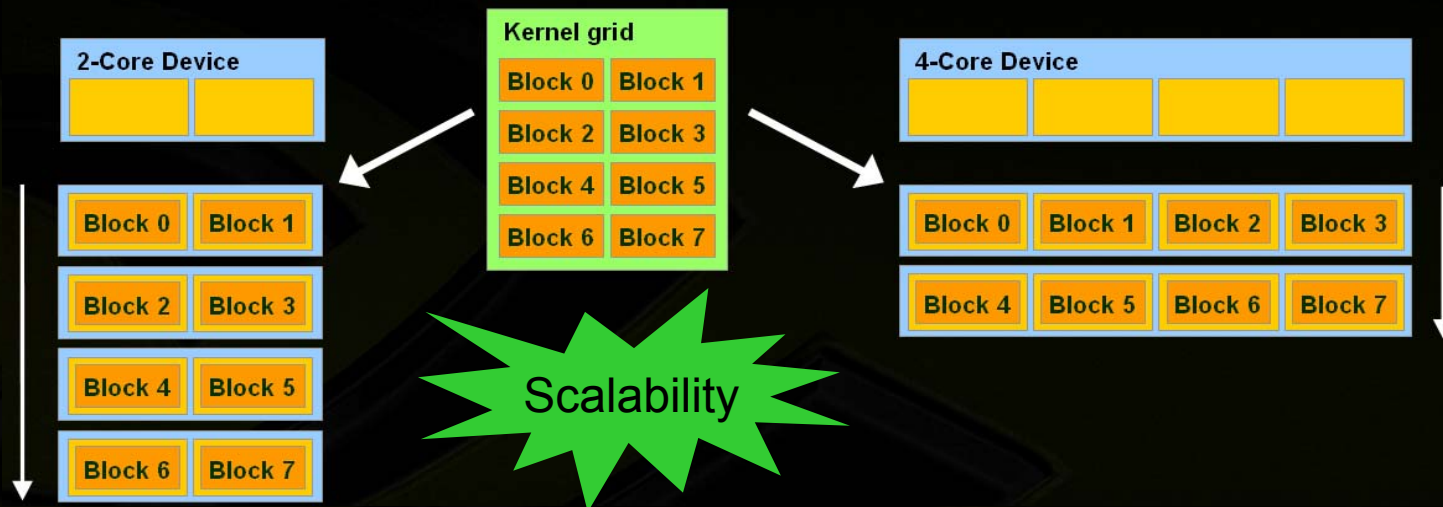
    return 0;
}
```

Blocks must be independent

- **Any possible interleaving of blocks should be valid**
 - presumed to run to completion without pre-emption
 - can run in any order
 - can run concurrently OR sequentially
- **Blocks may coordinate but not synchronize**
 - shared queue pointer: **OK**
 - shared lock: **BAD** can easily deadlock
- **Independence requirement gives scalability**

Blocks must be independent

- **Thread blocks can run in any order**
 - Concurrently or sequentially
 - Facilitates scaling of the same code across many devices





NVIDIA®

Coordinating CPU and GPU Execution



Synchronizing GPU and CPU

- **All kernel launches are asynchronous**
 - control returns to CPU immediately
 - kernel starts executing once all previous CUDA calls have completed
- **Memcopies are synchronous**
 - control returns to CPU once the copy is complete
 - copy starts once all previous CUDA calls have completed
- **cudaThreadSynchronize()**
 - blocks until all previous CUDA calls complete
- **Asynchronous CUDA calls provide:**
 - non-blocking memcopies
 - ability to overlap memcopies and kernel execution



CUDA Error Reporting to CPU

- **All CUDA calls return error code:**
 - except kernel launches
 - `cudaError_t` type
- **`cudaError_t cudaGetLastError(void)`**
 - returns the code for the last error (“no error” has a code)
- **`char* cudaGetErrorString(cudaError_t code)`**
 - returns a null-terminated character string describing the error

```
printf(“%s\n”, cudaGetErrorString( cudaGetLastError() ) );
```



CUDA Event API

- Events are inserted (recorded) into CUDA call streams
- Usage scenarios:
 - measure elapsed time for CUDA calls (clock cycle precision)
 - query the status of an asynchronous CUDA call
 - block CPU until CUDA calls prior to the event are completed
 - **asyncAPI** sample in CUDA SDK

```
cudaEvent_t start, stop;  
cudaEventCreate(&start);          cudaEventCreate(&stop);  
cudaEventRecord(start, 0);  
kernel<<<grid, block>>>(...);  
cudaEventRecord(stop, 0);  
cudaEventSynchronize(stop);  
float et;  
cudaEventElapsedTime(&et, start, stop);  
cudaEventDestroy(start);          cudaEventDestroy(stop);
```



Device Management

● CPU can query and select GPU devices

- `cudaGetDeviceCount(int* count)`
- `cudaSetDevice(int device)`
- `cudaGetDevice(int *current_device)`
- `cudaGetDeviceProperties(cudaDeviceProp* prop, int device)`
- `cudaChooseDevice(int *device, cudaDeviceProp* prop)`

● Multi-GPU setup:

- device 0 is used by default
- one CPU thread can control one GPU
 - multiple CPU threads can control the same GPU
 - calls are serialized by the driver



NVIDIA®

Shared Memory

Shared Memory



- **On-chip memory**

- 2 orders of magnitude lower latency than global memory
- Order of magnitude higher bandwidth than gmem
- **16KB** per multiprocessor
 - NVIDIA GPUs contain up to **30** multiprocessors

- **Allocated per threadblock**

- **Accessible by any thread in the threadblock**

- Not accessible to other threadblocks

- **Several uses:**

- Sharing data among threads in a threadblock
- User-managed cache (reducing gmem accesses)

Example of Using Shared Memory




- **Applying a 1D stencil:**

- 1D data

- For each output element, sum all elements within a radius

- **For example, radius = 3**


- Add 7 input elements



Implementation with Shared Memory



- 1D threadblocks (partition the output)
- Each threadblock outputs **BLOCK_DIMX** elements
 - Read input from gmem to smem
 - Needs **BLOCK_DIMX + 2*RADIUS** input elements
 - Compute
 - Write output to gmem



Kernel code

00056789101112
000567891011 6789101112000



```
__global__ void stencil( int *output, int *input, int dimx, int dimy )
{
    __shared__ int s_a[BLOCK_DIMX+2*RADIUS];

    int global_ix = blockIdx.x*blockDim.x + threadIdx.x;
    int local_ix  = threadIdx.x + RADIUS;

    s_a[local_ix] = input[global_ix];

    if ( threadIdx.x < RADIUS )
    {
        s_a[local_ix - RADIUS] = input[global_ix - RADIUS];
        s_a[local_ix + BLOCK_DIMX + RADIUS] = input[global_ix + RADIUS];
    }
    __syncthreads();

    int value = 0;
    for( offset = -RADIUS; offset<=RADIUS; offset++ )
        value += s_a[ local_ix + offset ];

    output[global_ix] = value;
}
```

sshould be gone

Thread Synchronization Function

- `void __syncthreads();`
- Synchronizes all threads in a thread-block
 - Since threads are scheduled at run-time
 - Once all threads have reached this point, execution resumes normally
 - Used to avoid RAW / WAR / WAW hazards when accessing shared memory
- Should be used in conditional code only if the conditional is uniform across the entire thread block

Memory Model Review

● Local storage

- Each thread has own local storage
- Mostly registers (managed by the compiler)
- Data lifetime = thread lifetime

● Shared memory

- Each thread block has own shared memory
 - Accessible only by threads within that block
- Data lifetime = block lifetime

● Global (device) memory

- Accessible by all threads as well as host (CPU)
- Data lifetime = from allocation to deallocation

Memory Model Review




Thread




Per-thread
Local Storage

Block





Per-block
Shared
Memory

Memory Model Review



Memory Model Review





NVIDIA®

CUDA Development Resources

CUDA Programming Resources



- **CUDA toolkit**

- Compiler, libraries, and documentation
- free download for Windows, Linux, and MacOS

- **CUDA SDK**

- code samples
- whitepapers

- **Instructional materials on CUDA Zone**

- slides and audio
- parallel programming course at University of Illinois UC
- tutorials
- forums

GPU Tools



● Profiler

- Available now for all supported OSs
- Command-line or GUI
- Sampling signals on GPU for:
 - Memory access parameters
 - Execution (serialization, divergence)

● Debugger

- Currently linux only (gdb)
- Runs on the GPU

● Emulation mode